


7.2


How Convolutional Neural Networks Work

Part 2: Convolutional Layers


Sebastian Raschka and the Lightning AI Team


Looking at **convolutional layers** in more detail






Let's zoom into this part







Input (image)




Feature map



Input (image)





Feature map




Input (image)

3x3 feature detector (kernel, filter)




Feature map

Slide feature detector over inputs




Input (image)




Feature map

Slide feature detector over inputs




Input (image)




Feature map

Slide feature detector over inputs

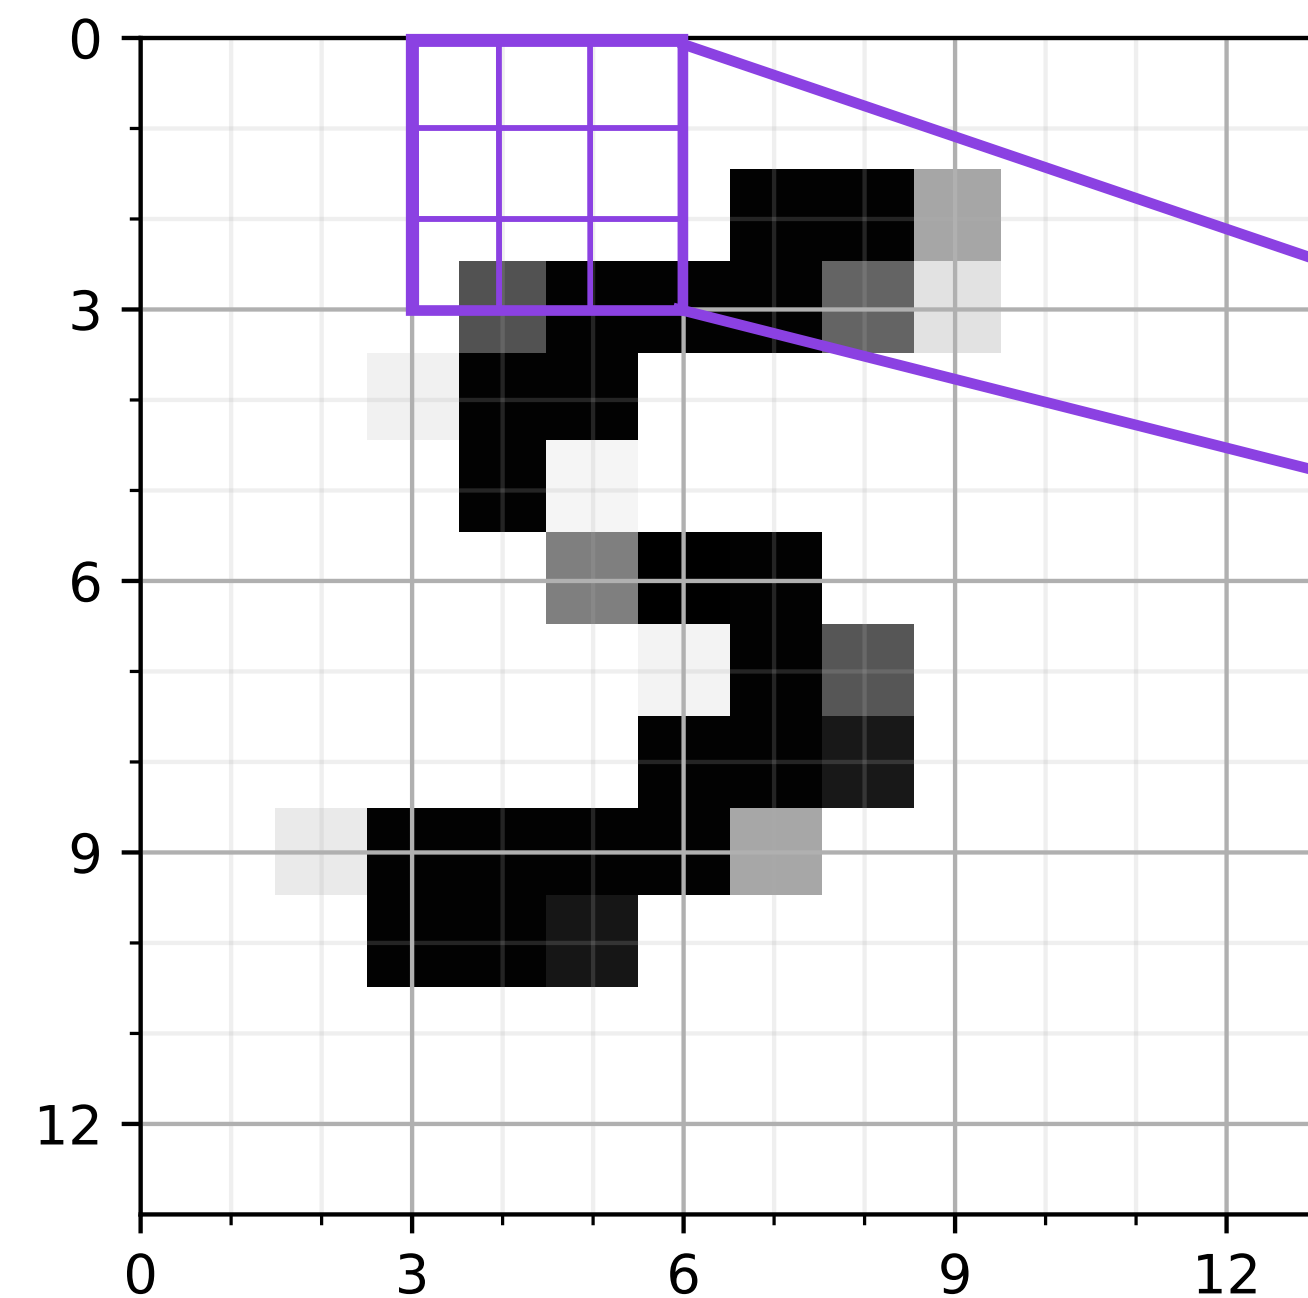


Input (image)




Feature map

Slide feature detector over inputs




Input (image)




Feature map


Slide feature detector over inputs



Input (image)




Feature map




Input (image)

Slide feature detector over inputs




Feature map

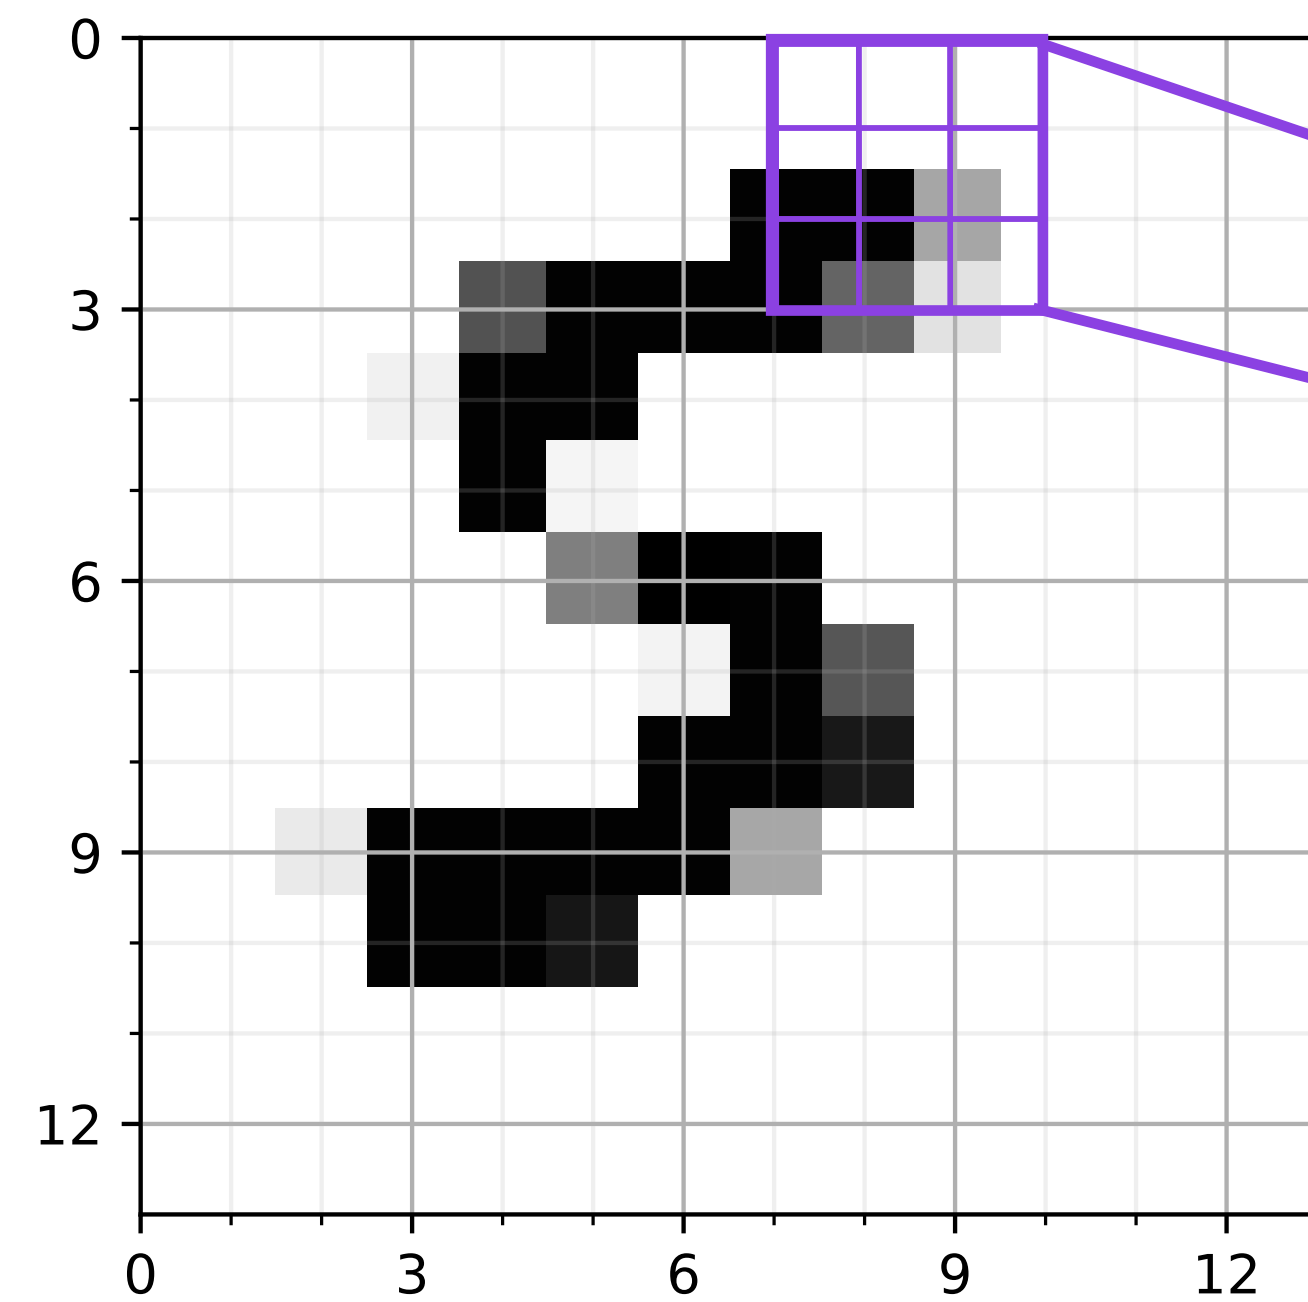
Slide feature detector over inputs



Input (image)




Feature map




Input (image)

Slide feature detector over inputs




Feature map




Input (image)

Slide feature detector over inputs




Feature map




Input (image)

Slide feature detector over inputs




Feature map

Slide feature detector over inputs




Input (image)




Feature map

Slide feature detector over inputs




Input (image)




Feature map


What is happening during this operation?




What is happening during this operation?



What is happening during this operation?




$$z = b + \sum_j w_j x_j$$




The inputs (x 's) differ as we slide over the image.

w_1	w_2	w_3
w_4	w_5	w_6
w_7	w_8	w_9

The weights (w 's) do not differ.




$$z = b + \sum_j w_j x_j$$




The inputs (x 's) differ as we slide over the image.

w_1	w_2	w_3
w_4	w_5	w_6
w_7	w_8	w_9


The weights (w 's) do not differ. → weight sharing



$$z = b + \sum_j w_j x_j$$



w_1	w_2	w_3
w_4	w_5	w_6
w_7	w_8	w_9



weight sharing


- Rationale: A feature detector that works well in one region may also work well in another region
- A reduction in parameters to fit (compared to MLP)

Convolutional layer


w_1	w_2	w_3
w_4	w_5	w_6
w_7	w_8	w_9

b

Consists of weights



Convolutional layer



Consists of weights


```
layer = torch.nn.Conv2d(1, 1, kernel_size=3)
```

```
layer.weight
```

Parameter containing:

```
tensor([[[[ 0.0930,  0.1602,  0.1177],  
          [-0.0802, -0.0701, -0.2747],  
          [ 0.1806,  0.2647,  0.2281]]]], requires_grad=True)
```

Convolutional layer



Consists of weights & bias unit(s)

```
layer = torch.nn.Conv2d(1, 1, kernel_size=3)
```

```
layer.weight
```

Parameter containing:

```
tensor([[[[ 0.0930,  0.1602,  0.1177],  
          [-0.0802, -0.0701, -0.2747],  
          [ 0.1806,  0.2647,  0.2281]]]], requires_grad=True)
```

```
layer.bias
```

Parameter containing:

```
tensor([-0.2351], requires_grad=True)
```

Of course, we are free to choose different kernel sizes

```
layer = torch.nn.Conv2d(1, 1, kernel_size=3)
```


```
layer.weight
```

```
Parameter containing:
tensor([[[[ 0.0930,  0.1602,  0.1177],
           [-0.0802, -0.0701, -0.2747],
           [ 0.1806,  0.2647,  0.2281]]]], requires_grad=True)
```


```
layer.bias
```

```
Parameter containing:
tensor([-0.2351], requires_grad=True)
```

So far, we looked at an excerpt



So far, we looked at an excerpt



What about the other channels?

NEXT: let's learn about convolutions with multiple channels

