7.2

How Convolutional Neural Networks Work

Part 3: Convolutions With Multiple Channels

Sebastian Raschka and the Lightning AI Team

We looked at convolutional layers in more detail

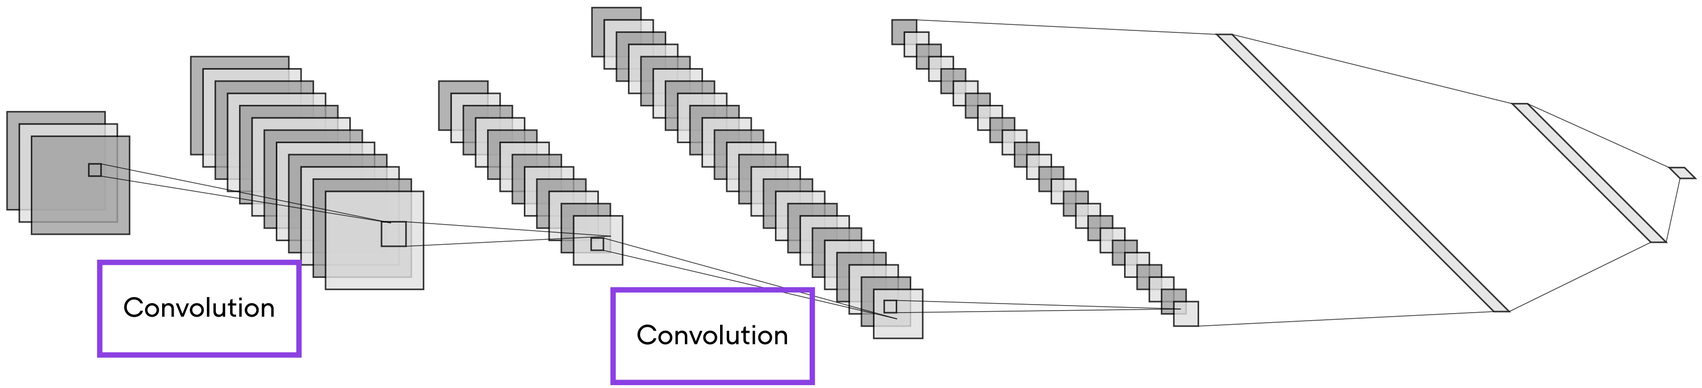

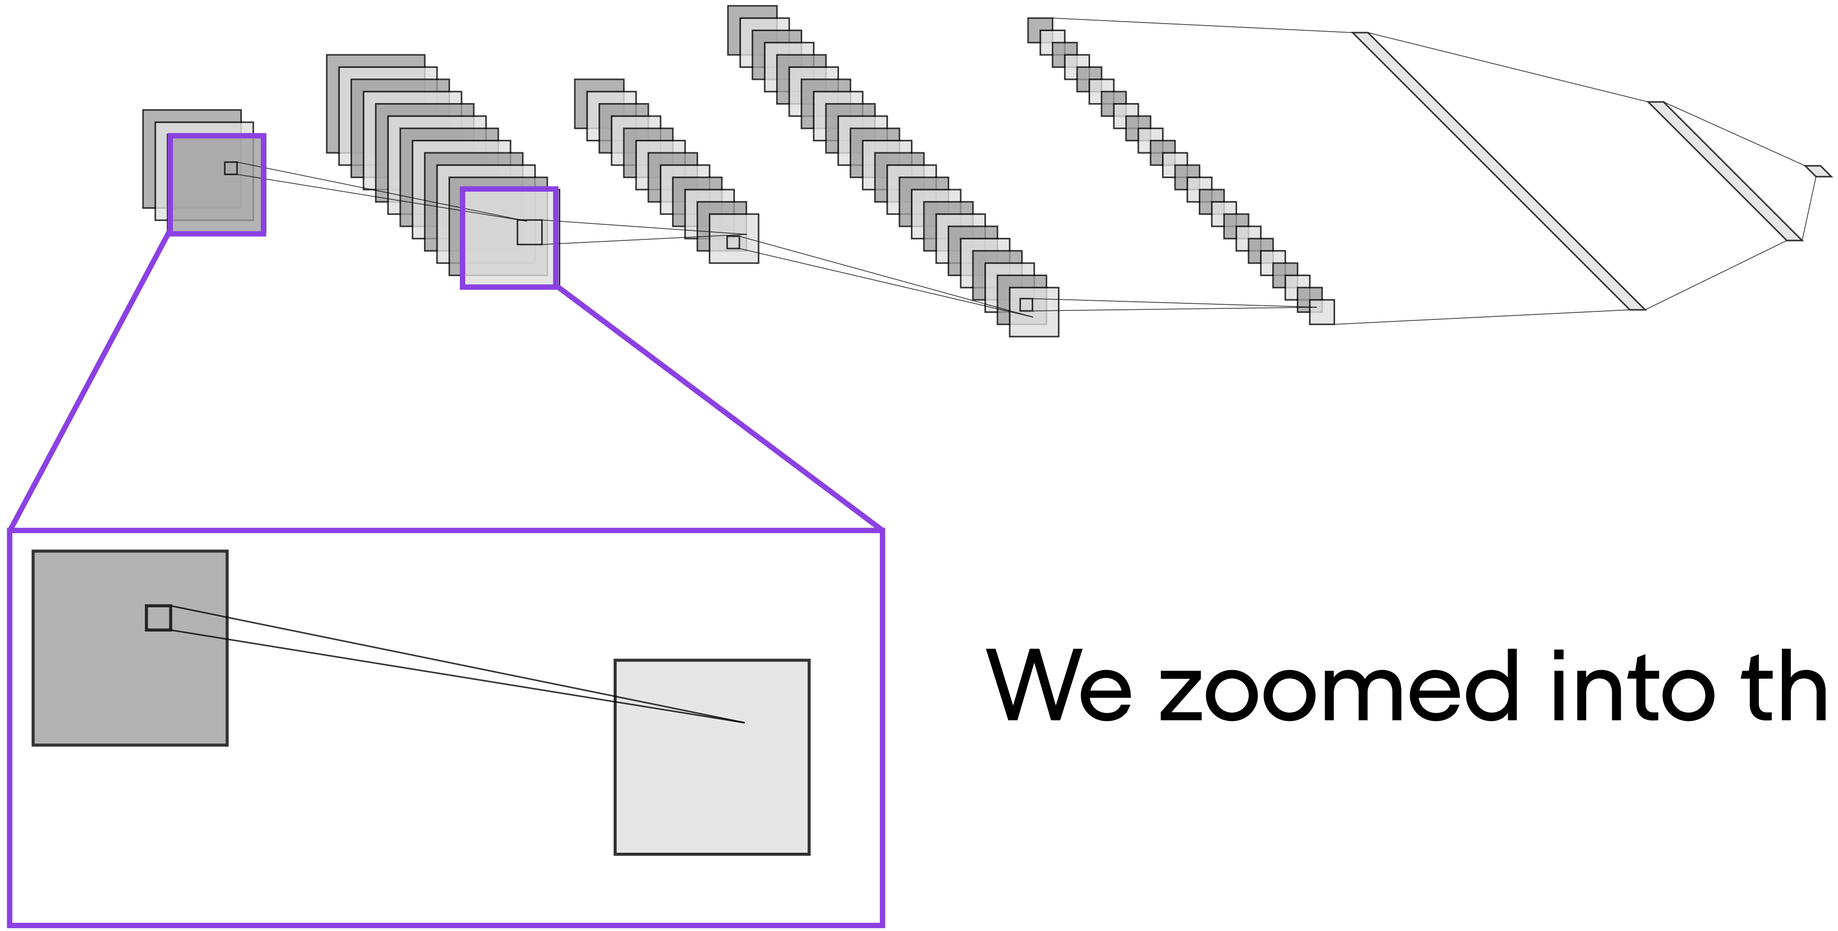

We zoomed into this part

1 input channel, 1 output channel

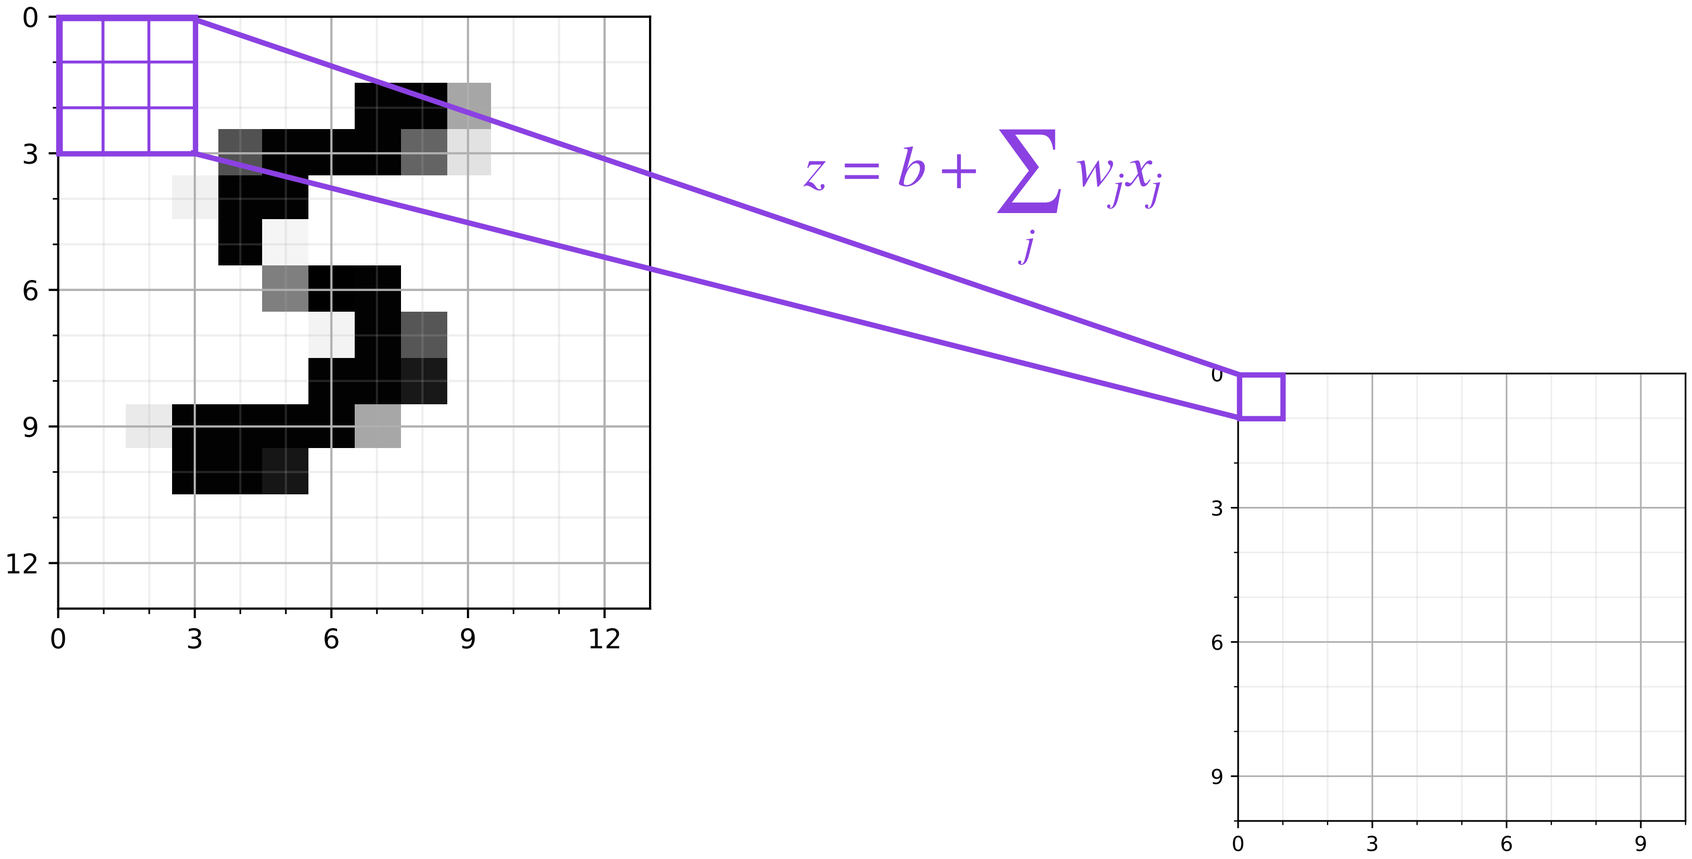

Now, let's learn about convolutions with multiple channels

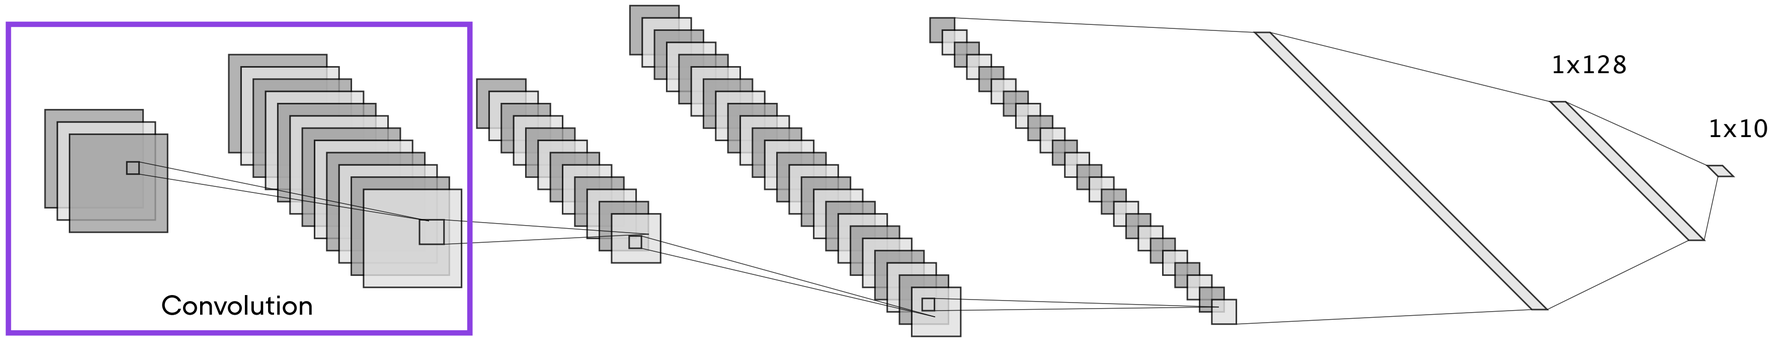

Let's start with multiple output channels

1 input channel, 3 output channels

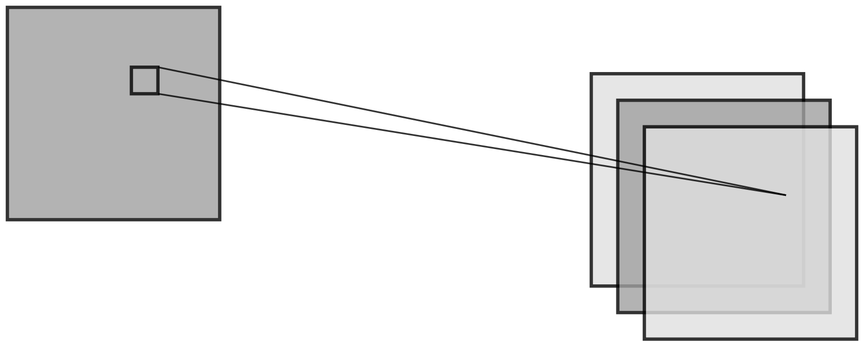

1 input channel, 3 output channels

Adding a bit more annotation ...

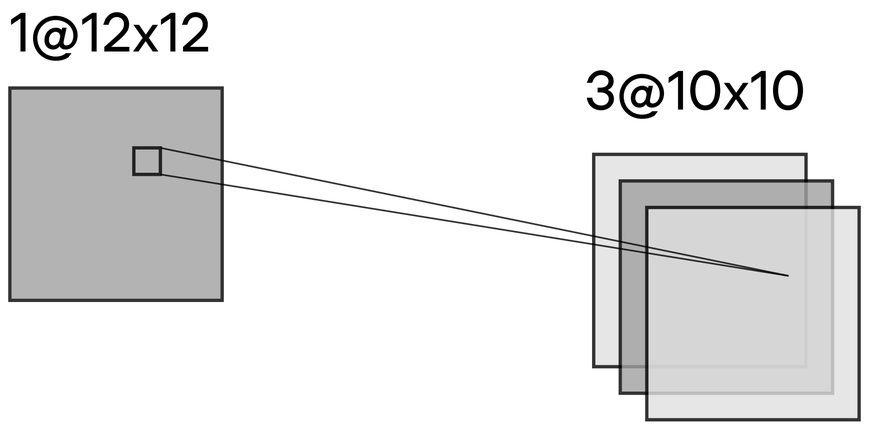

We will discuss how to control the output size in a later video

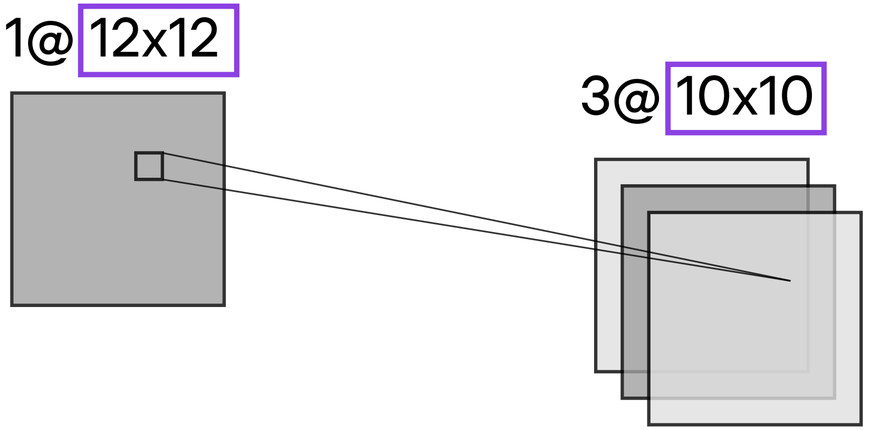

1 input channel, 3 output channels

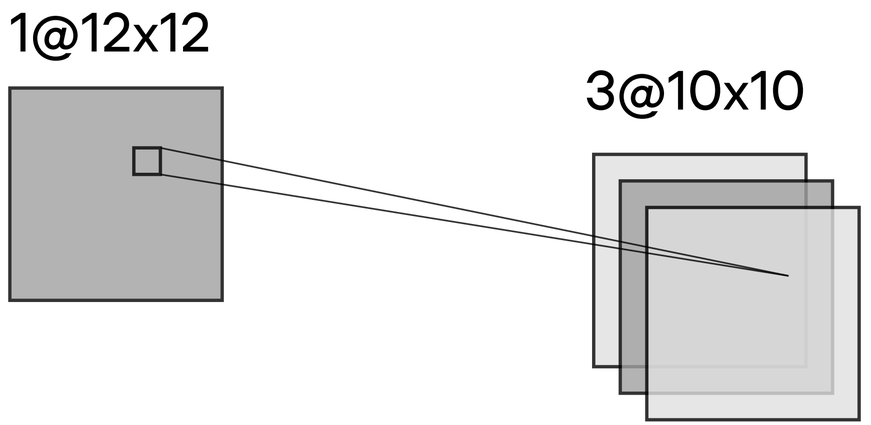

1 input channel, 3 output channels

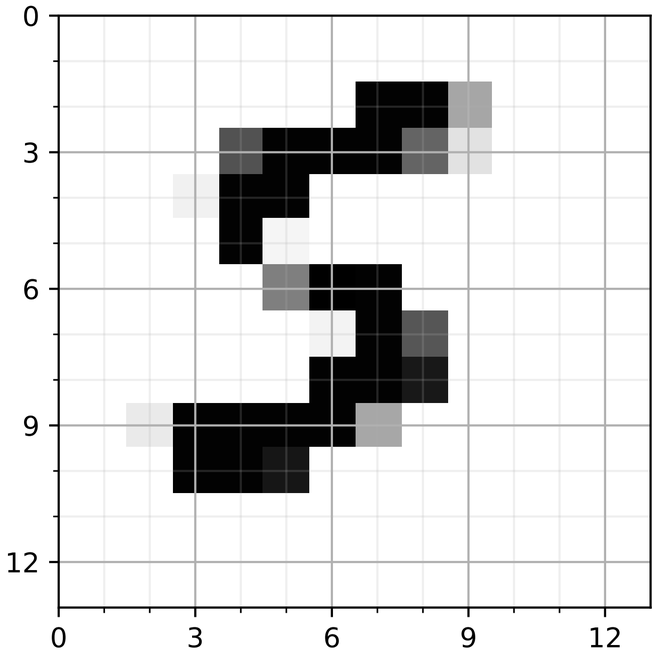

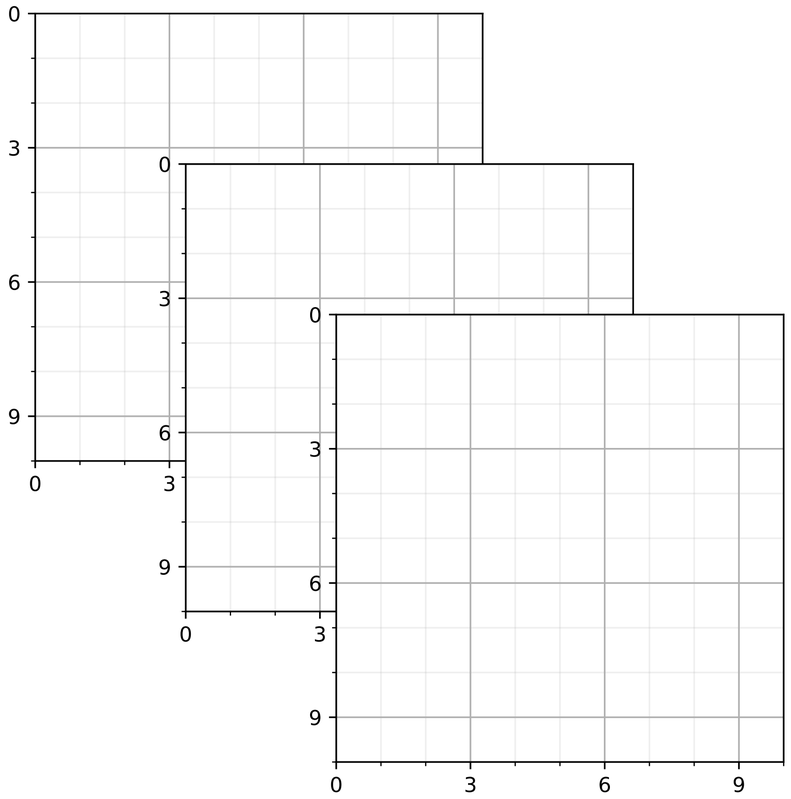

1 input channel, 3 output channels

$$z = b^{(@1)} + \sum_j w_j^{(@1)} x_j$$

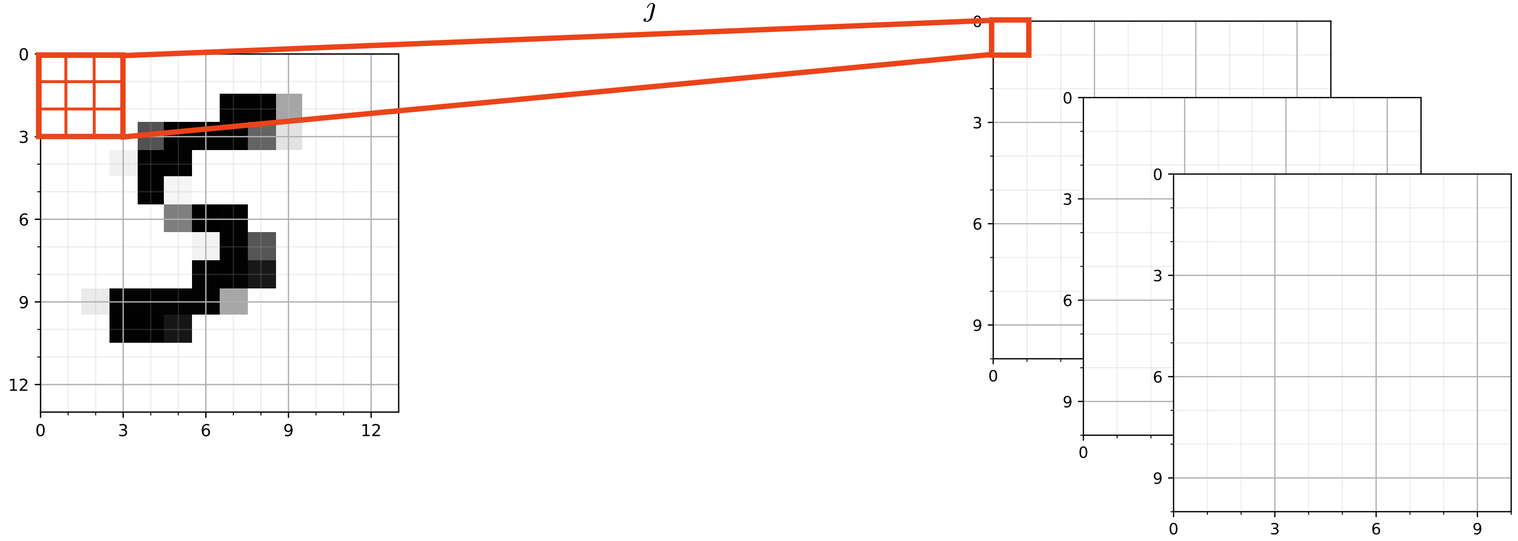

1 input channel, 3 output channels

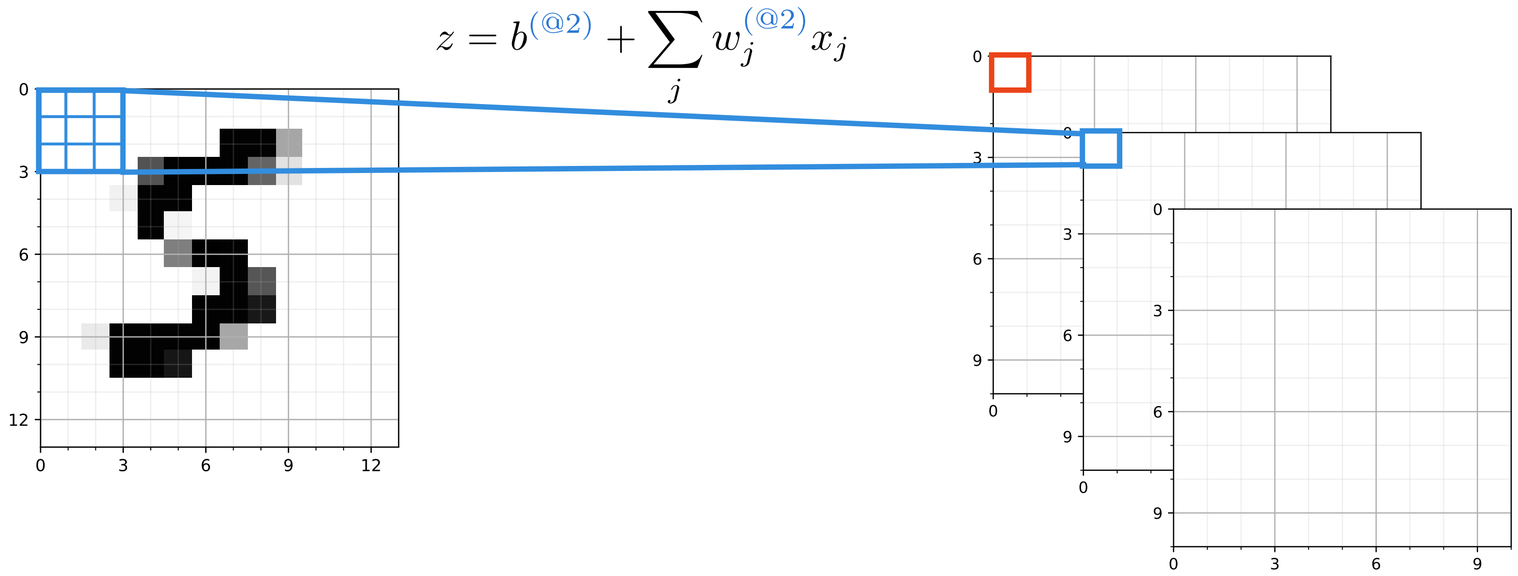

1 input channel, 3 output channels

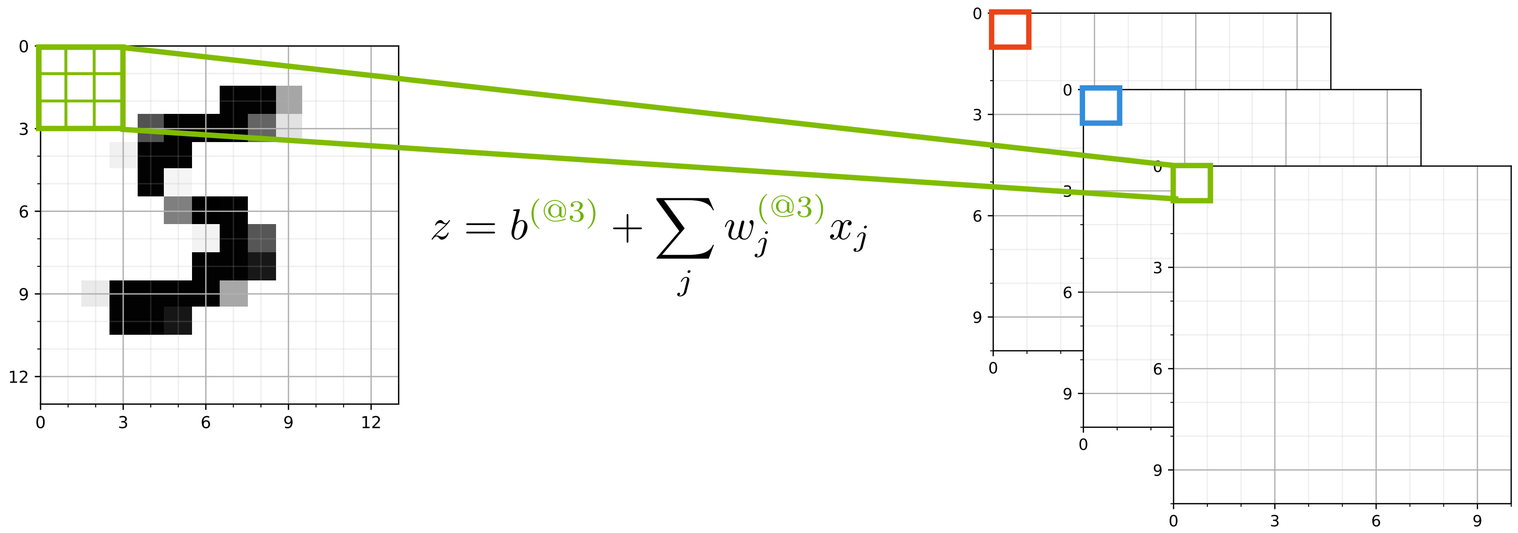

1 input channel, 3 output channels

1@12x12

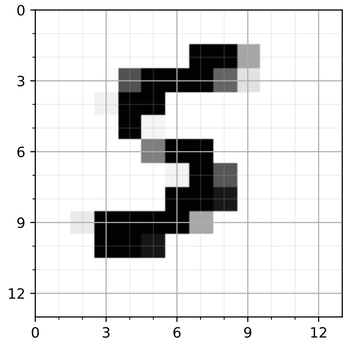

3@10x10

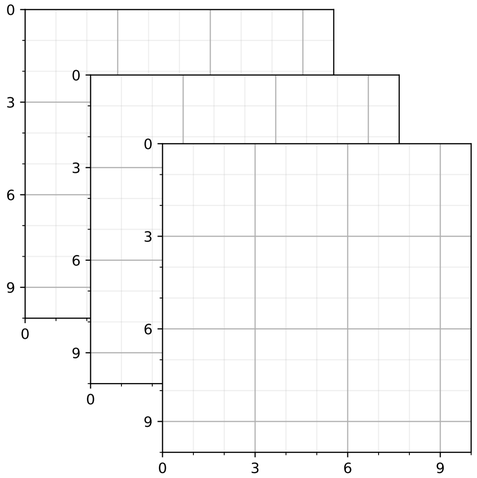

Multiple "feature detectors" (kernels) are used to create multiple feature maps

What about multiple input channels?

RGB Image

R

G

B

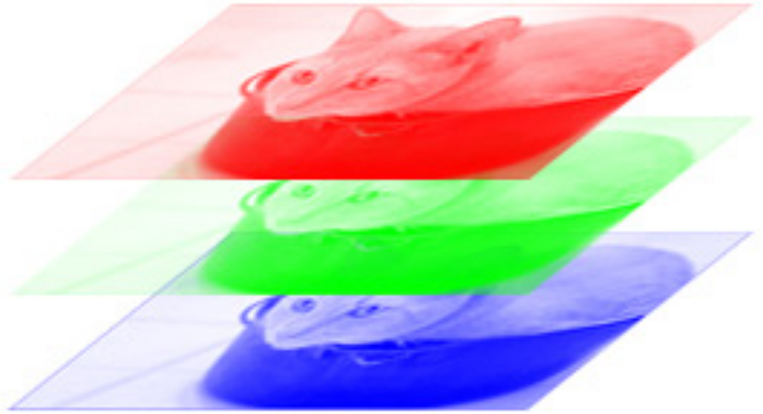

=

Image Source: https://code.tutsplus.com/tutorials/create-a-retro-crt-distortion-effect-using-rgb-shifting--active-3359

Color image as a stack of matrices

3 input channels, 1 output channel

kernel has 3 channels

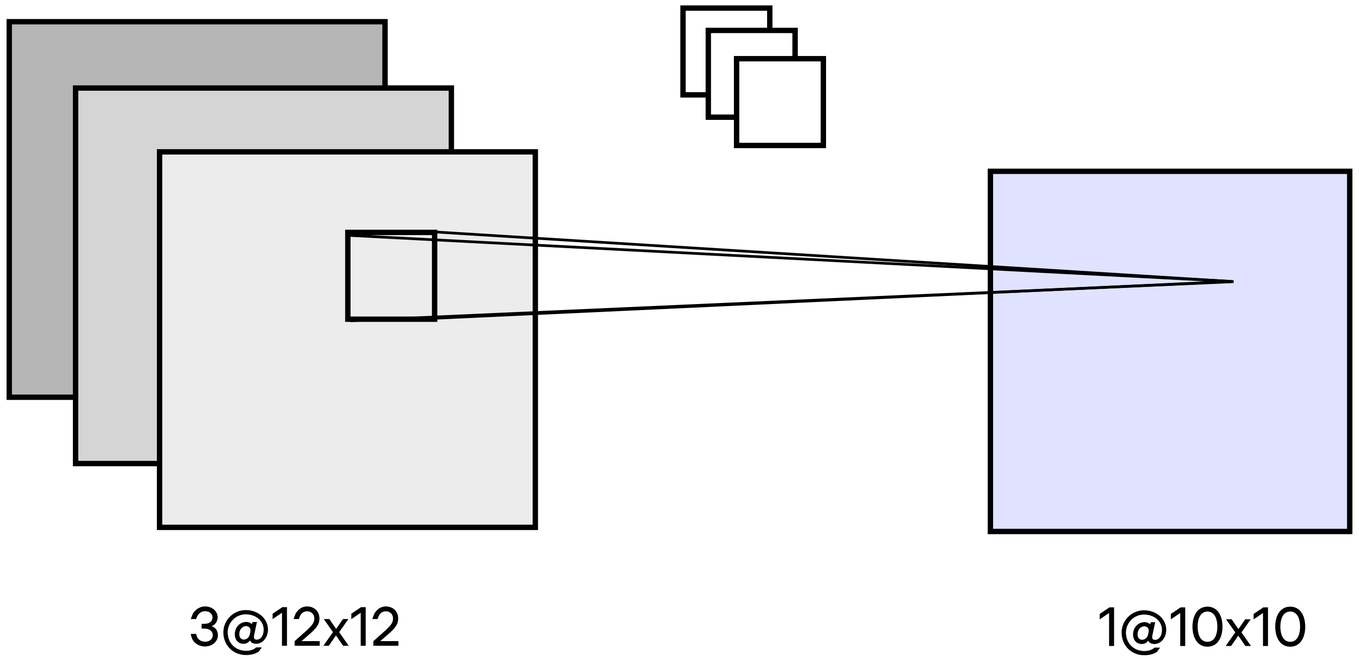

3@12x12

1@10x10

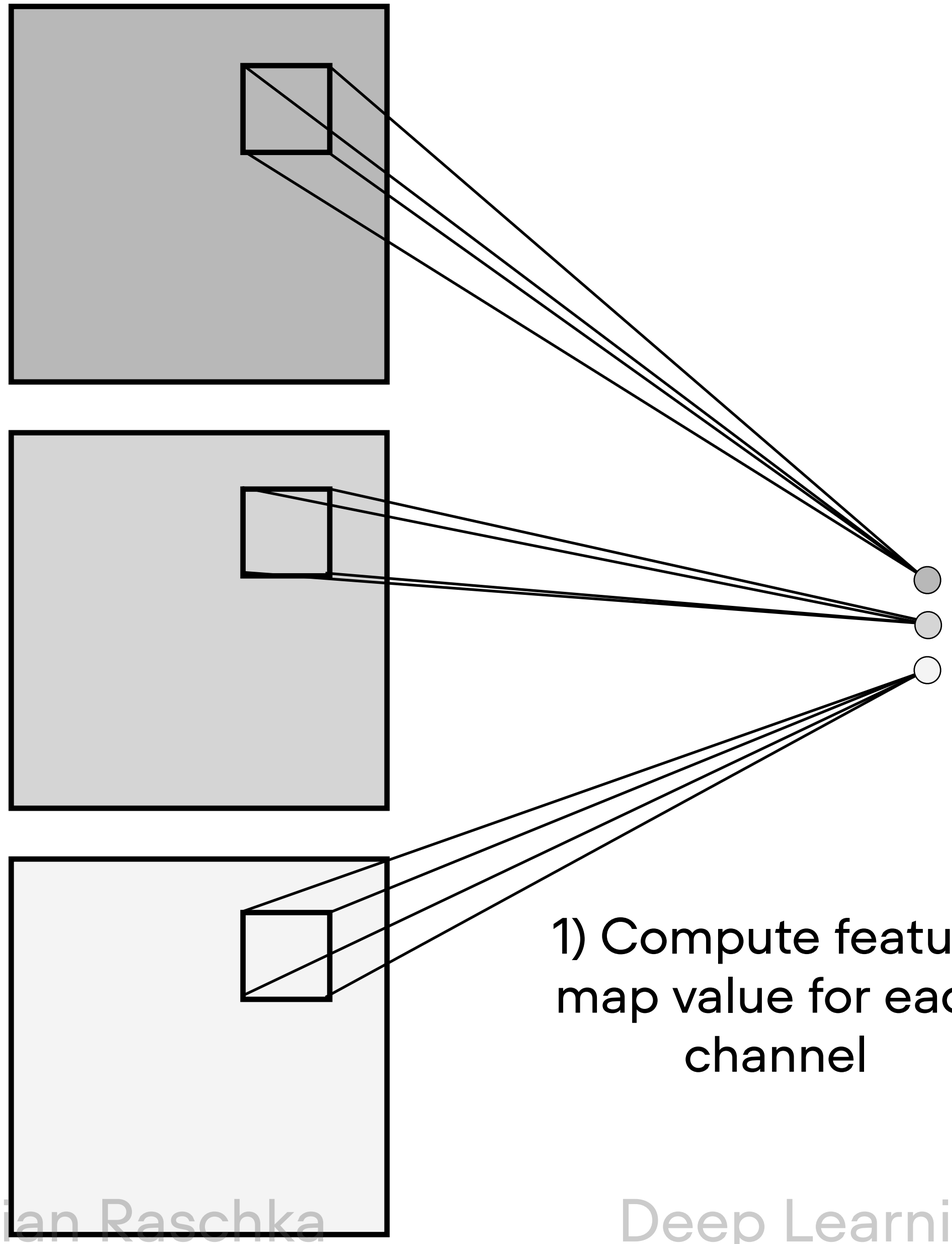

3@12x12

1) Compute feature map value for each channel

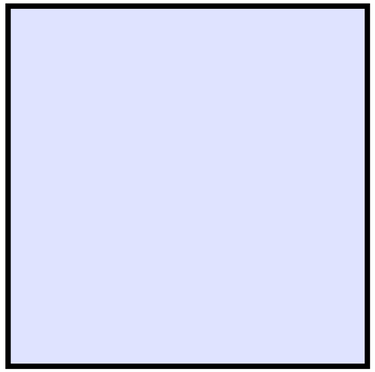

1@10x10

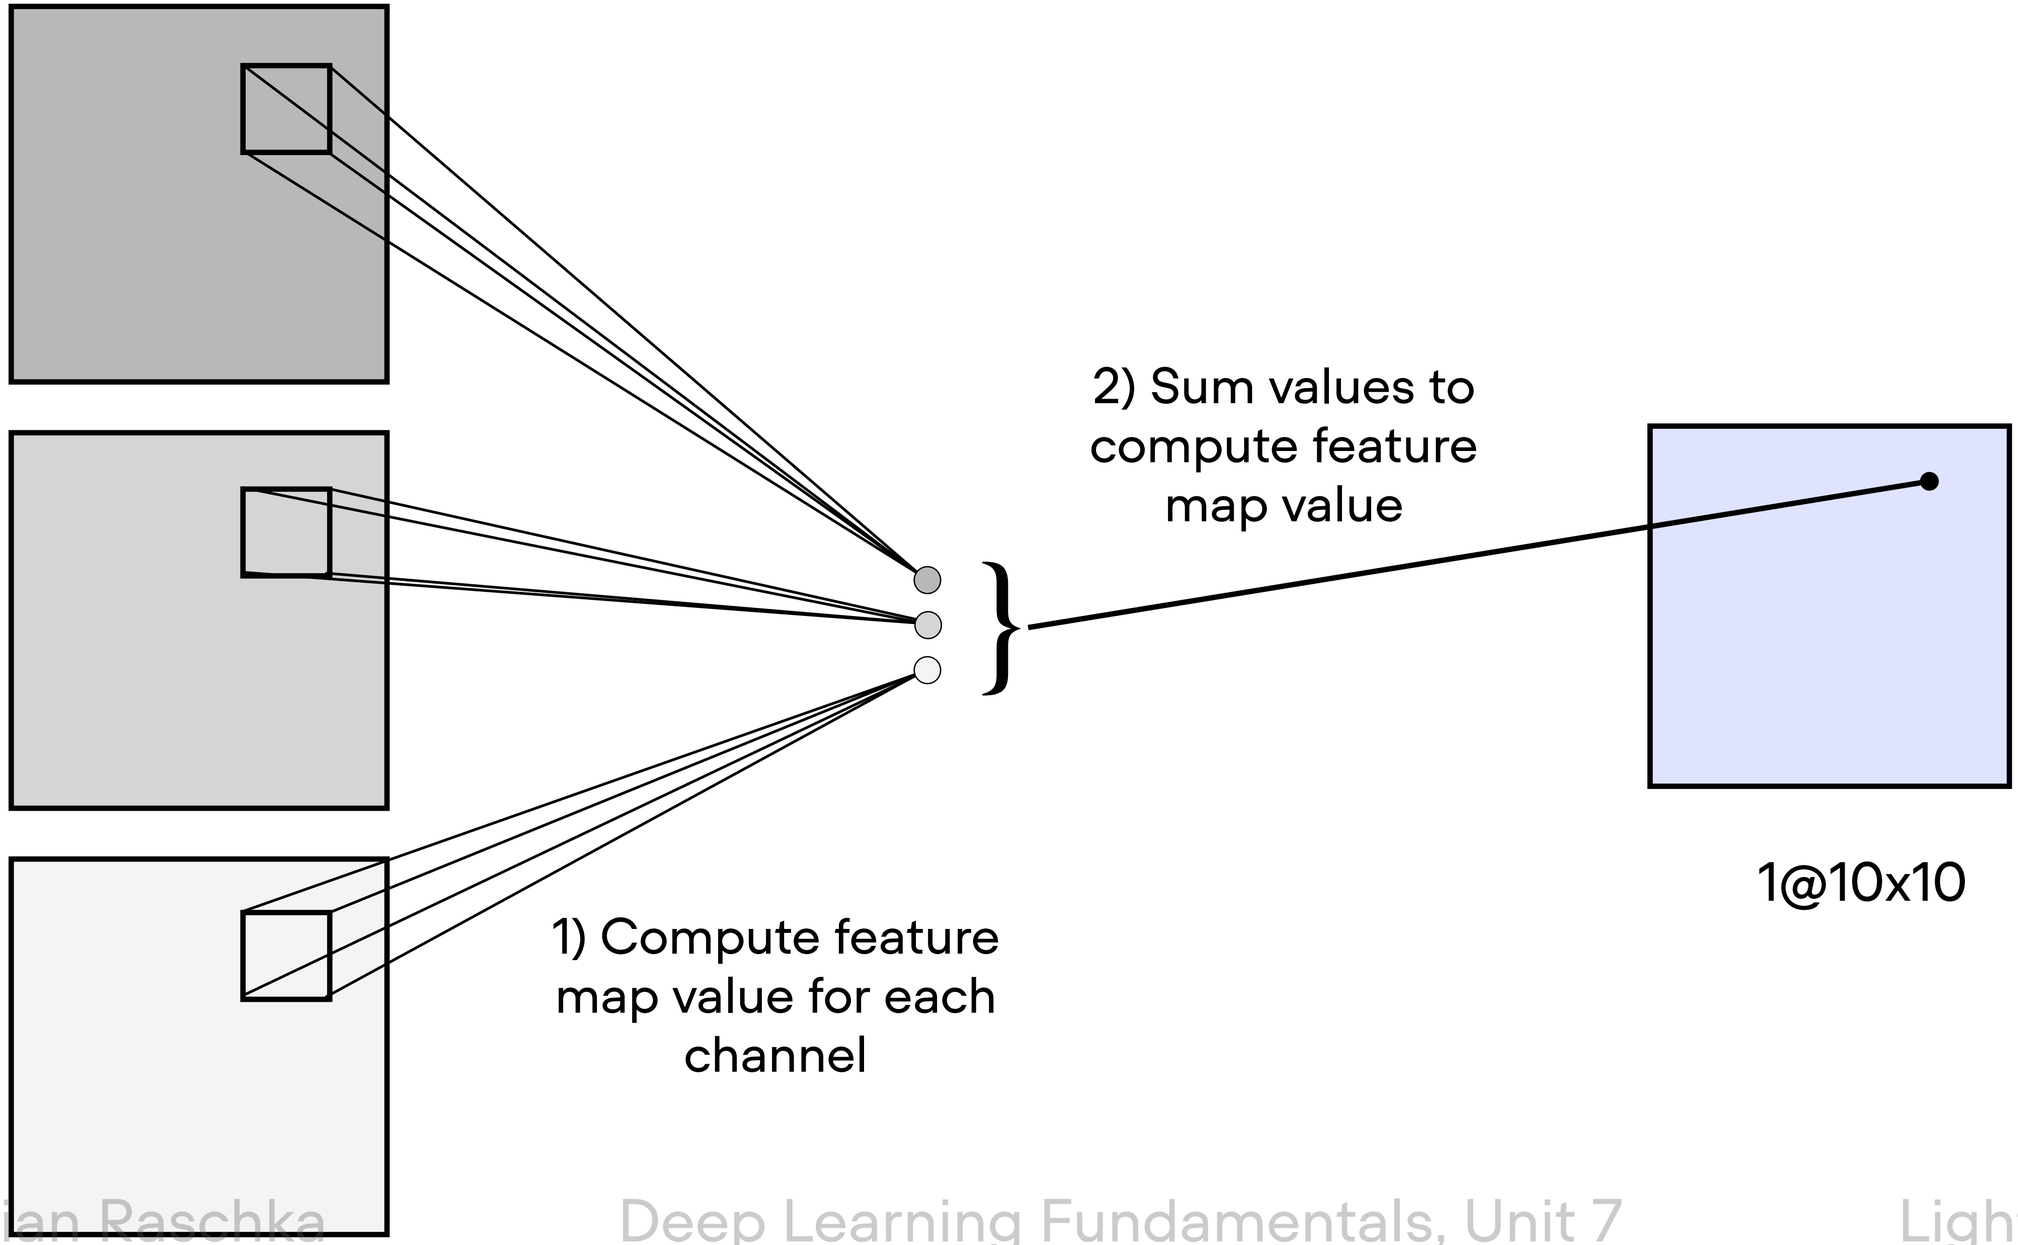
2) Sum values to compute feature map value
1@10x10
1) Compute feature map value for each channel
3@12x12

What about multiple input AND output channels?

3 input channels

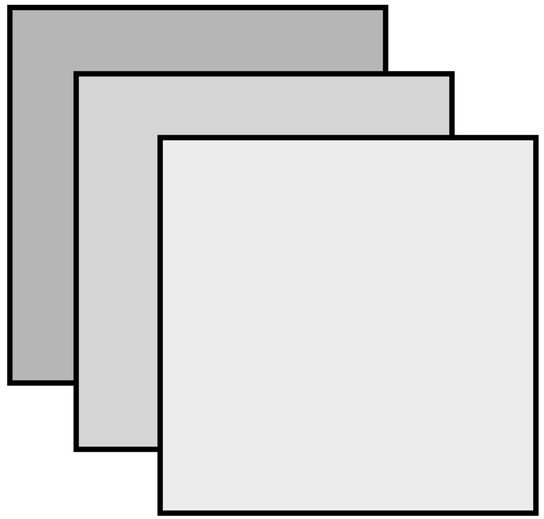

3@64x64

5 output channels

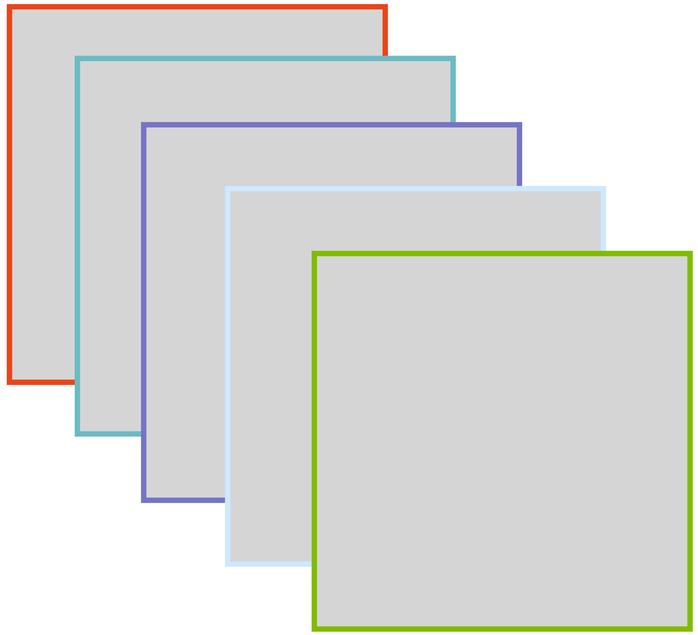

5@64x64

3 input channels

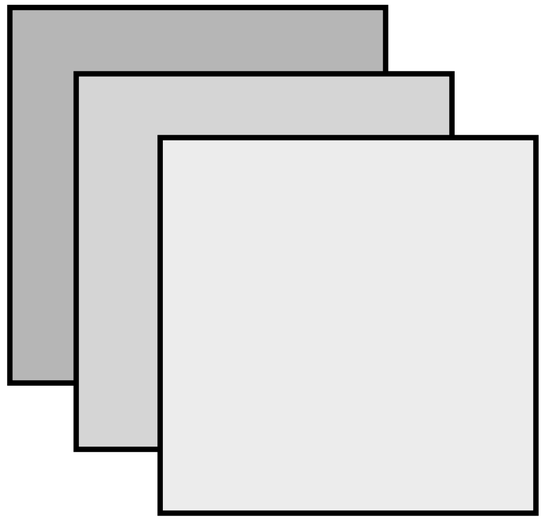

3@64x64

5 output channels

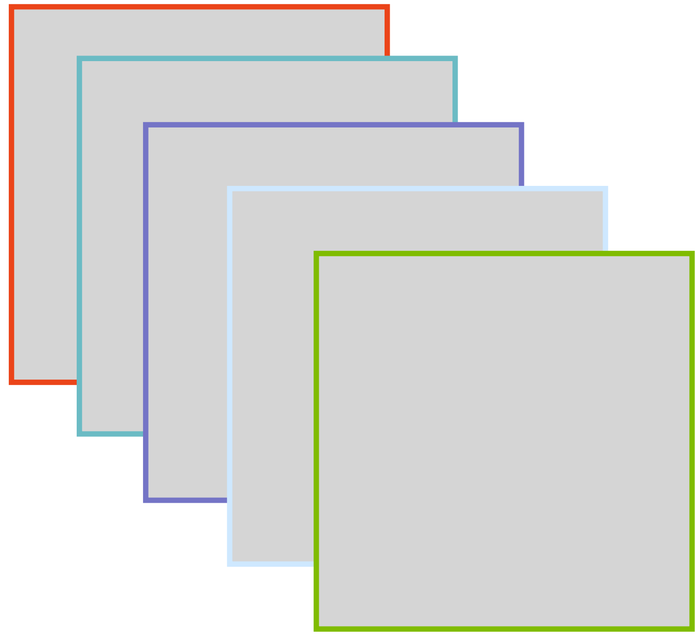

5@64x64

5 kernels with
3 channels each

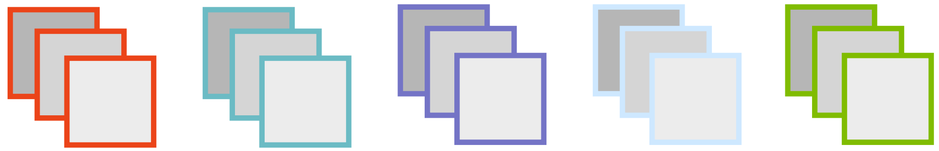

```python
import torch


layer = torch.nn.Conv2d(in_channels=3, out_channels=5, kernel_size=2)
layer.weight.shape
```

```
torch.Size([5, 3, 2, 2])
```

5 kernels with
3 channels each

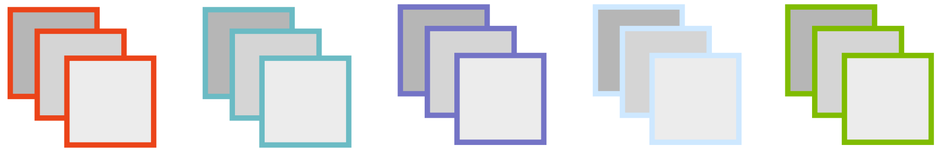

We covered convolutional layers with multiple channels

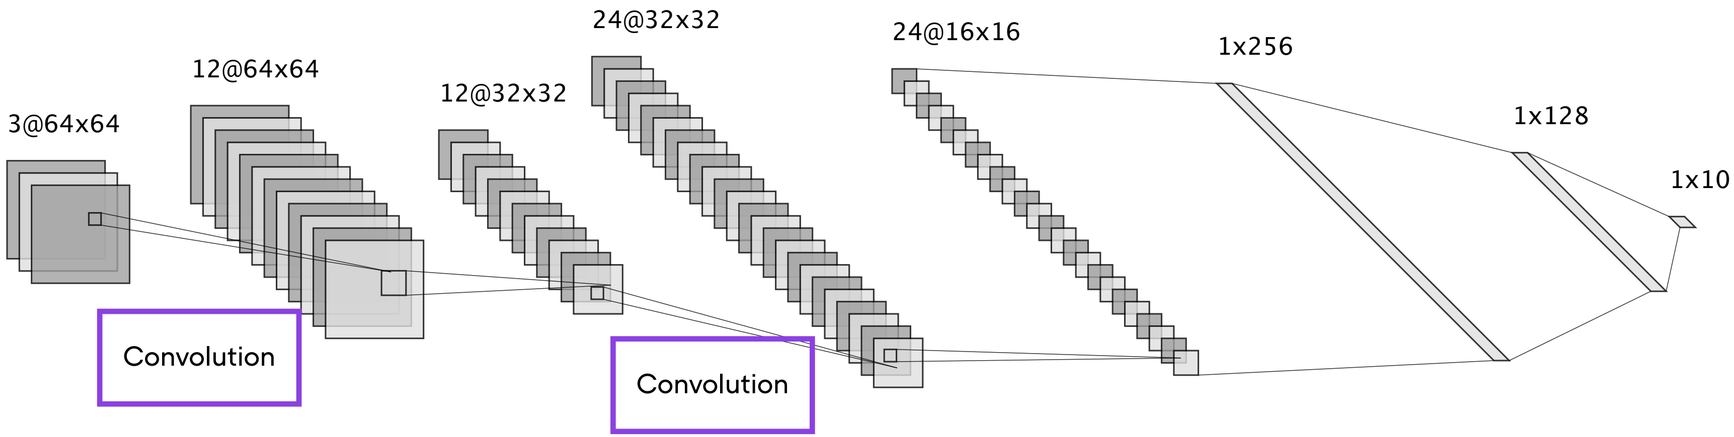

Next: What are these pooling layers?