


# 7.2

## How Convolutional Neural Networks Work


### Part 4: What Are Pooling Layers?

Sebastian Raschka and the Lightning AI Team

# Let's talk about pooling layers



# CNN architectures are like a funnel



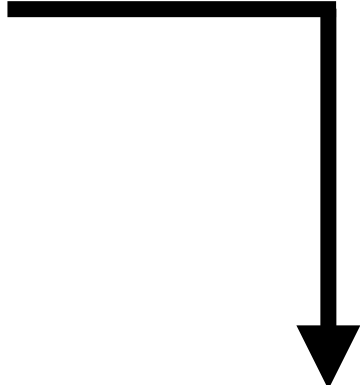
# CNN architectures are like a funnel

Squeezing out information,  
learning representations

# CNN architectures are like a funnel

We decrease height and width  
and increase the number of channels

We decrease height and width  
and increase the number of channels

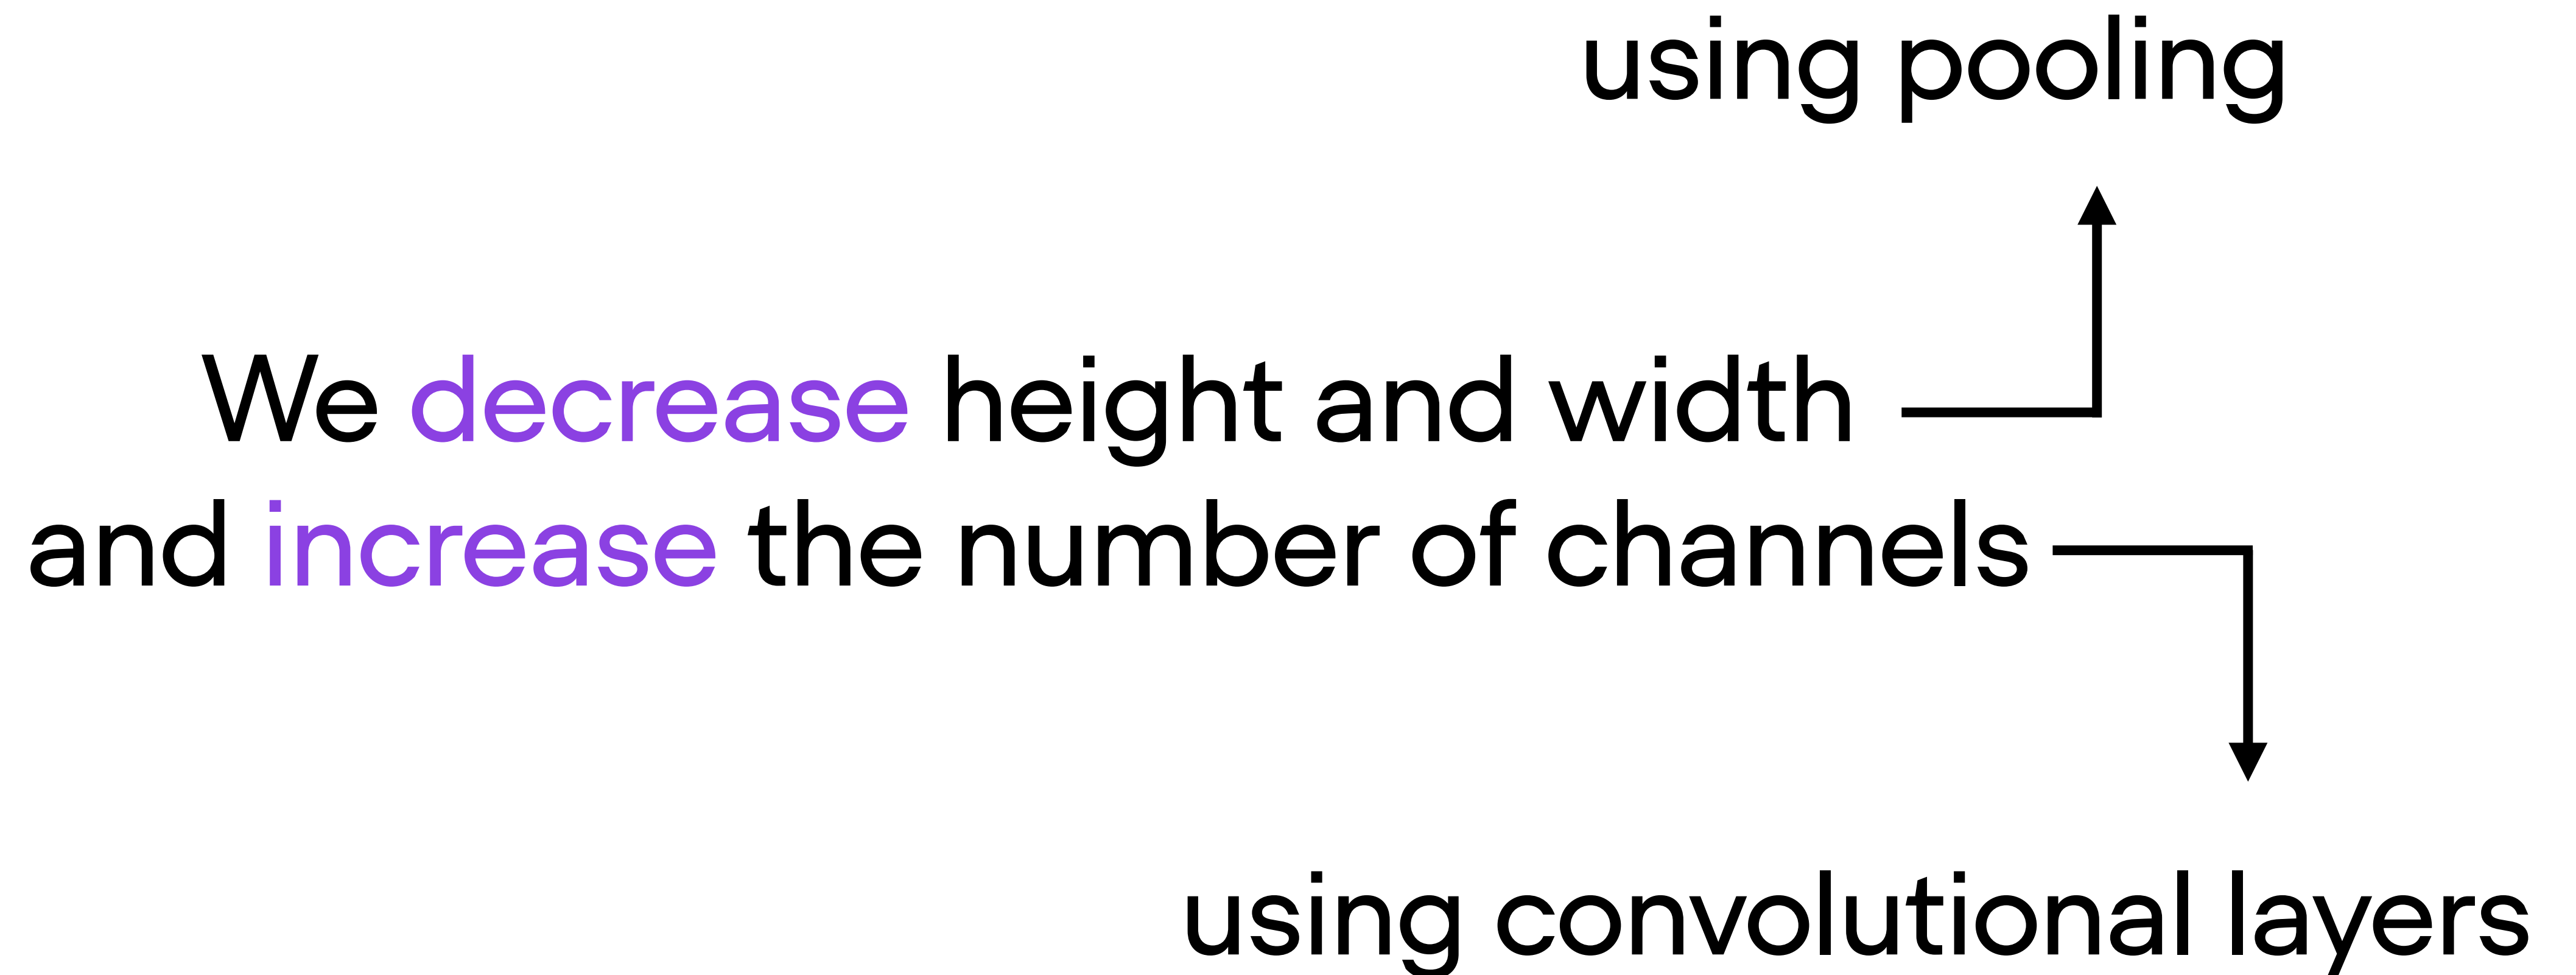


using convolutional layers

using pooling

We decrease height and width  
and increase the number of channels

using convolutional layers




# Max Pooling

with the  
following input:

1	-4	2	4	-2	1
-5	5	7	3	-4	9
2	-3	3	-2	2	1
-1	3	8	1	8	7
9	2	-4	3	1	-9
8	3	8	2	3	1




# Max Pooling



and kernel size  
2x2

1	-4	2	4	-2	1
-5	5	7	3	-4	9
2	-3	3	-2	2	1
-1	3	8	1	8	7
9	2	-4	3	1	-9
8	3	8	2	3	1


# Max Pooling



input

output


# Max Pooling



input

output


# Max Pooling



input

output


# Max Pooling



input

output


# Max Pooling



input

output


# Max Pooling



input

output

# Max Pooling




input

output




# Max Pooling



input

output

# Max Pooling




input

output

# Average Pooling

# Average Pooling



input

output

Typical pooling layers don't have any learnable parameters

Pooling layers can help with local invariance

... but some information will be lost

Some architectures skip pooling layers altogether

Springenberg, Dosovitskiy, Brox, Riedmiller (2014). Striving for Simplicity: The All Convolutional Net. <https://arxiv.org/abs/1412.6806>

We can reduce the feature map size  
with more "stride"



## A typical setup

```
import torch
import torch.nn as nn

layers_with_pooling = nn.Sequential(
    nn.Conv2d(3, 8, kernel_size=3),
    nn.MaxPool2d(kernel_size=2, stride=2),
    nn.Conv2d(8, 16, kernel_size=3),
    nn.MaxPool2d(kernel_size=2, stride=2),
)

example = torch.rand(3, 110, 110)
layers_with_pooling(example).shape

torch.Size([16, 26, 26])
```

## A typical setup

```
import torch
import torch.nn as nn

layers_with_pooling = nn.Sequential(
    nn.Conv2d(3, 8, kernel_size=3),
    nn.MaxPool2d(kernel_size=2, stride=2),
    nn.Conv2d(8, 16, kernel_size=3),
    nn.MaxPool2d(kernel_size=2, stride=2),
)

example = torch.rand(3, 110, 110)
layers_with_pooling(example).shape
```

```
torch.Size([16, 26, 26])
```

## without pooling layers

```
layers_no_pooling = nn.Sequential(
    nn.Conv2d(3, 8, kernel_size=3),
    #nn.MaxPool2d(kernel_size=2, stride=2),
    nn.Conv2d(8, 16, kernel_size=3),
    #nn.MaxPool2d(kernel_size=2, stride=2),
)

example = torch.rand(3, 110, 110)
layers_no_pooling(example).shape
```

```
torch.Size([16, 106, 106])
```

## A typical setup

```
import torch
import torch.nn as nn

layers_with_pooling = nn.Sequential(
    nn.Conv2d(3, 8, kernel_size=3),
    nn.MaxPool2d(kernel_size=2, stride=2),
    nn.Conv2d(8, 16, kernel_size=3),
    nn.MaxPool2d(kernel_size=2, stride=2),
)

example = torch.rand(3, 110, 110)
layers_with_pooling(example).shape

torch.Size([16, 26, 26])
```

without pooling layers  
but with stride=2

```
import torch
import torch.nn as nn

layers_no_pooling_2 = nn.Sequential(
    nn.Conv2d(3, 8, kernel_size=3, stride=2),
    #nn.MaxPool2d(kernel_size=2, stride=2),
    nn.Conv2d(8, 16, kernel_size=3, stride=2),
    #nn.MaxPool2d(kernel_size=2, stride=2),
)

example = torch.rand(3, 110, 110)
layers_no_pooling_2(example).shape

torch.Size([16, 26, 26])
```

Next: More control over the output size via padding