

7.3

Convolutional Neural Network Architectures

Part 1: The Main Ideas

Sebastian Raschka and the Lightning AI Team

Summary of the main concepts behind CNNs

- **Sparse-connectivity:** A single element in the feature map is connected to only a small patch of pixels (versus fully-connected layers in MLPs)

Summary of the main concepts behind CNNs

- **Sparse-connectivity:** A single element in the feature map is connected to only a small patch of pixels (versus fully-connected layers in MLPs)
- **Parameter-sharing:** The same weights are used for different patches of the input image.

Summary of the main concepts behind CNNs

- **Sparse-connectivity:** A single element in the feature map is connected to only a small patch of pixels (versus fully-connected layers in MLPs)
- **Parameter-sharing:** The same weights are used for different patches of the input image.
- **Many layers:** Combining extracted local patterns to global patterns

- **Sparse-connectivity** \longrightarrow **Neighboring pixels are related**
- **Parameter-sharing** \longrightarrow **Same filters work in different parts of the image**

Summary of the main concepts behind CNNs

- **Sparse-connectivity** → **Neighboring pixels are related**
- **Parameter-sharing** → **Same filters work in different parts of the image**

"Inductive biases" that help CNNs learn more quickly and generalize better
(compared to fully-connected networks)

- **Sparse-connectivity**
 - **Parameter-sharing**
- } **Nice side effect: convolutional layers are small!**

```

class PyTorchMLP(torch.nn.Module):
    def __init__(self):
        super().__init__()

        self.all_layers = torch.nn.Sequential(
            # 1st hidden layer
            torch.nn.Linear(3*224*224, 10_000),
            torch.nn.ReLU(),

            # 2nd hidden layer
            torch.nn.Linear(10_000, 1_000),
            torch.nn.ReLU(),

            # 3rd hidden layer
            torch.nn.Linear(1_000, 100),
            torch.nn.ReLU(),

            # output layer
            torch.nn.Linear(100, 10),
        )

    def forward(self, x):
        x = torch.flatten(x, start_dim=1)
        logits = self.all_layers(x)
        return logits

```



```

class PyTorchMLP(torch.nn.Module):
    def __init__(self):
        super().__init__()

        self.all_layers = torch.nn.Sequential(
            # 1st hidden layer
            torch.nn.Linear(3*224*224, 10_000),
            torch.nn.ReLU(),

            # 2nd hidden layer
            torch.nn.Linear(10_000, 1_000),
            torch.nn.ReLU(),

            # 3rd hidden layer
            torch.nn.Linear(1_000, 100),
            torch.nn.ReLU(),

            # output layer
            torch.nn.Linear(100, 10),
        )

    def forward(self, x):
        x = torch.flatten(x, start_dim=1)
        logits = self.all_layers(x)
        return logits

```

$3 \times 224 \times 224 \times 10,000 + 10,000 = 1,505,290,000$

```

class PyTorchMLP(torch.nn.Module):
    def __init__(self):
        super().__init__()

        self.all_layers = torch.nn.Sequential(
            # 1st hidden layer
            torch.nn.Linear(3*224*224, 10_000), → 3*224*224 * 10,000 + 10,000 = 1,505,290,000
            torch.nn.ReLU(),

            # 2nd hidden layer
            torch.nn.Linear(10_000, 1_000), → 10,001,000
            torch.nn.ReLU(),

            # 3rd hidden layer
            torch.nn.Linear(1_000, 100), → 100,100
            torch.nn.ReLU(),

            # output layer
            torch.nn.Linear(100, 10), → 1,010
        )

    def forward(self, x):
        x = torch.flatten(x, start_dim=1)
        logits = self.all_layers(x)
        return logits

```

1,515,392,110 parameters!!!

```

class PyTorchMLP(torch.nn.Module):
    def __init__(self):
        super().__init__()

        self.all_layers = torch.nn.Sequential(
            # 1st hidden layer
            torch.nn.Linear(3*224*224, 10_000), —————→ 3*224*224 * 10,000 + 10,000 = 1,505,290,000
            torch.nn.ReLU(),

            # 2nd hidden layer
            torch.nn.Linear(10_000, 1_000), —————→ 10,001,000
            torch.nn.ReLU(),

            # 3rd hidden layer
            torch.nn.Linear(1_000, 100), —————→ 100,100
            torch.nn.ReLU(),

            # output layer
            torch.nn.Linear(100, 10), —————→ 1,010
        )

    def forward(self, x):
        x = torch.flatten(x, start_dim=1)
        logits = self.all_layers(x)
        return logits

```

1,515,392,110 parameters!!!

```

size = 0.
for name, param in mlp.named_parameters():
    size += sys.getsizeof(param.storage()) / 1_024**3
print(f"Model size: {size:.2f} GB")

```

Model size: 5.65 GB


```
Sequential(
  (0): Conv2d(3, 8, kernel_size=(5, 5), stride=(2, 2)) → 3*224*224 * 10,000 + 10,000 = 1,505,290,000
  (1): ReLU(inplace=True)
  (2): Conv2d(8, 24, kernel_size=(5, 5), stride=(2, 2)) → 4,824
  (3): ReLU(inplace=True)
  (4): Conv2d(24, 32, kernel_size=(3, 3), stride=(2, 2)) → 6,944
  (5): ReLU(inplace=True)
  (6): Conv2d(32, 48, kernel_size=(3, 3), stride=(2, 2)) → 13,872
  (7): ReLU(inplace=True)
  (8): Linear(48*12*12, 200) → 1,382,600
  (9): ReLU(inplace=True)
  (10): Linear(200, 10) → 2,010
)

1,515,392,110 parameters!!!
```

```
parameters():
  total() / 1_024**3
```

```
class PyTorchCNN(torch.nn.Module):
    def __init__(self):
        super().__init__()

        self.cnn_layers = torch.nn.Sequential(
            torch.nn.Conv2d(3, 8, kernel_size=5, stride=2), → 3*5*5*8 + 8 = 608
            torch.nn.ReLU(),

            torch.nn.Conv2d(8, 24, kernel_size=5, stride=2), → 4,824
            torch.nn.ReLU(),

            torch.nn.Conv2d(24, 32, kernel_size=3, stride=2), → 6,944
            torch.nn.ReLU(),

            torch.nn.Conv2d(32, 48, kernel_size=3, stride=2), → 13,872
            torch.nn.ReLU(),
        )

        self.fc_layers = torch.nn.Sequential(
            torch.nn.Linear(48*12*12, 200), → 1,382,600
            torch.nn.ReLU(),
            torch.nn.Linear(200, 10), → 2,010
        )

    def forward(self, x):

        x = self.cnn_layers(x)
        x = torch.flatten(x, start_dim=1)
        logits = self.fc_layers(x)
        return logits

1,410,858 parameters
```

```
size = 0.
for name, param in cnn.named_parameters():
    size += sys.getsizeof(param.storage()) / 1_024**3
print(f"Model size: {size:.3f} GB")
Model size: 0.016 GB
```

Next: CNNs And Their Inception(s)