


8.3

Introduction to Recurrent Neural Networks

Part 3: Encoding Inputs Using Embedding Layers

Sebastian Raschka and the Lightning AI Team

Recurrent neural networks (RNNS) for modeling sequences



Recurrent neural network (RNN)

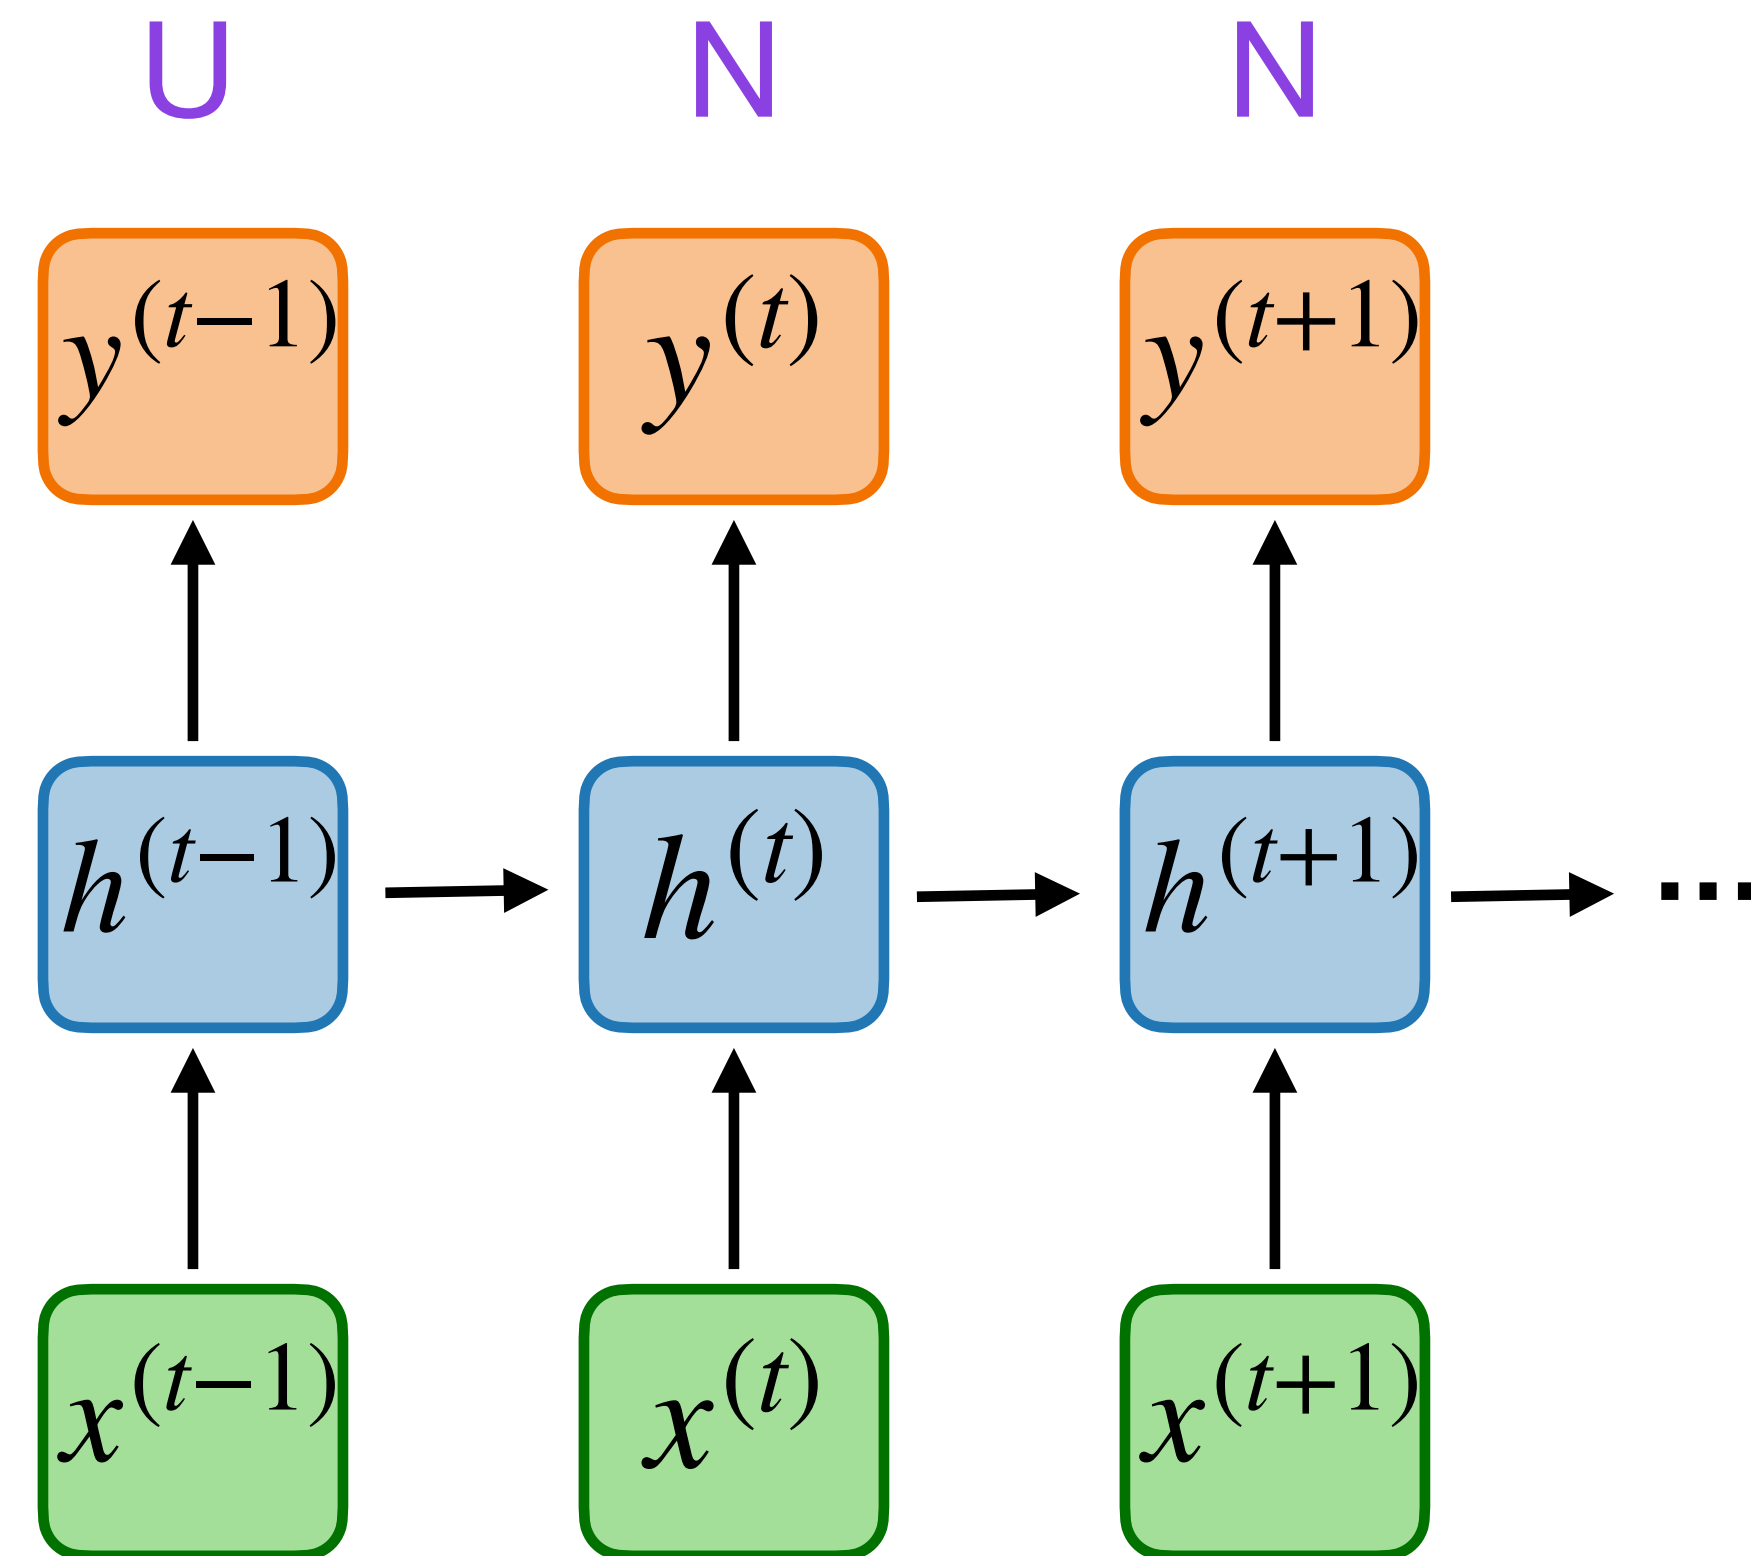
Deep Learning Fundamentals, Unit 3 The same RNN Lightning AI

**Suppose we implement a simple RNN that predicts
the next character in a sentence:**

“Sunny days are the best days to go for a walk or have a picnic.”

A simple RNN that predicts the next character

Desired outputs:



Sentence:

"Sunny days are the best days to go for a walk or have a picnic."

Character inputs:

S

U

N

Remember one-hot encoding?

	A	B	C	D	E	F	G	H	I	J	K	L	M	N	O	P	Q	R	S	T	U	V	W	X	Y	Z
S	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	1	0	0	0	0	0	0	0
U	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	1	0	0	0	0	0
N	0	0	0	0	0	0	0	0	0	0	0	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0
N	0	0	0	0	0	0	0	0	0	0	0	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0
Y	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	1	0

One-hot encoded input letter

S

A	B	C	D	E	F	G	H	I	J	K	L	M	N	O	P	Q	R	S	T	U	V	W	X	Y	Z
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	1	0	0	0	0	0	0	0



[0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 1., 0., 0., 0., 0., 0., 0., 0.]

One-hot encoded input letter

S

A	B	C	D	E	F	G	H	I	J	K	L	M	N	O	P	Q	R	S	T	U	V	W	X	Y	Z
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	1	0	0	0	0	0	0	0

[0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 1., 0., 0., 0., 0., 0., 0., 0.]

Weight matrix

[[0.6912,	0.8765,	0.4939],
		0.6342,	0.7481,	0.7717],
		0.8395,	0.2128,	0.3696],
		0.4900,	0.1509,	0.0689],
		0.2587,	0.9171,	0.8670],
		0.7213,	0.9922,	0.5701],
		0.7598,	0.5231,	0.3666],
		0.5150,	0.5216,	0.9682],
		0.2248,	0.0261,	0.4427],
		0.1818,	0.6863,	0.8713],
		0.4192,	0.1566,	0.9004],
		0.8102,	0.5741,	0.4241],
		0.1116,	0.0466,	0.2786],
		0.9816,	0.7363,	0.5899],
		0.9224,	0.3672,	0.6972],
		0.1207,	0.3372,	0.2128],
		0.0660,	0.1524,	0.8440],
		0.2162,	0.5640,	0.0988],
		0.2605,	0.3766,	0.3502],
		0.2334,	0.4757,	0.7581],
		0.7382,	0.9807,	0.4762],
		0.2369,	0.8102,	0.8798],
		0.6932,	0.2671,	0.8018],
		0.9593,	0.5302,	0.4290],
		0.6231,	0.8825,	0.8836],
		0.4623,	0.8503,	0.7279]]

Multiply input vector with weight matrix

```
[0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 1., 0., 0., 0., 0., 0., 0.]
```

```
[[0.6912, 0.8765, 0.4939],  
 [0.6342, 0.7481, 0.7717],  
 [0.8395, 0.2128, 0.3696],  
 [0.4900, 0.1509, 0.0689],  
 [0.2587, 0.9171, 0.8670],  
 [0.7213, 0.9922, 0.5701],  
 [0.7598, 0.5231, 0.3666],  
 [0.5150, 0.5216, 0.9682],  
 [0.2248, 0.0261, 0.4427],  
 [0.1818, 0.6863, 0.8713],  
 [0.4192, 0.1566, 0.9004],  
 [0.8102, 0.5741, 0.4241],  
 [0.1116, 0.0466, 0.2786],  
 [0.9816, 0.7363, 0.5899],  
 [0.9224, 0.3672, 0.6972],  
 [0.1207, 0.3372, 0.2128],  
 [0.0660, 0.1524, 0.8440],  
 [0.2162, 0.5640, 0.0988],  
 [0.2605, 0.3766, 0.3502],  
 [0.2334, 0.4757, 0.7581],  
 [0.7382, 0.9807, 0.4762],  
 [0.2369, 0.8102, 0.8798],  
 [0.6932, 0.2671, 0.8018],  
 [0.9593, 0.5302, 0.4290],  
 [0.6231, 0.8825, 0.8836],  
 [0.4623, 0.8503, 0.7279]]
```


Multiply input vector with weight matrix

$$0 \times 0.6912 = 0$$

[0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 1., 0., 0., 0., 0., 0., 0.]

[0.6912,	0.8765,	0.4939],
[0.6342,	0.7481,	0.7717],
[0.8395,	0.2128,	0.3696],
[0.4900,	0.1509,	0.0689],
[0.2587,	0.9171,	0.8670],
[0.7213,	0.9922,	0.5701],
[0.7598,	0.5231,	0.3666],
[0.5150,	0.5216,	0.9682],
[0.2248,	0.0261,	0.4427],
[0.1818,	0.6863,	0.8713],
[0.4192,	0.1566,	0.9004],
[0.8102,	0.5741,	0.4241],
[0.1116,	0.0466,	0.2786],
[0.9816,	0.7363,	0.5899],
[0.9224,	0.3672,	0.6972],
[0.1207,	0.3372,	0.2128],
[0.0660,	0.1524,	0.8440],
[0.2162,	0.5640,	0.0988],
[0.2605,	0.3766,	0.3502],
[0.2334,	0.4757,	0.7581],
[0.7382,	0.9807,	0.4762],
[0.2369,	0.8102,	0.8798],
[0.6932,	0.2671,	0.8018],
[0.9593,	0.5302,	0.4290],
[0.6231,	0.8825,	0.8836],
[0.4623,	0.8503,	0.7279]

Multiply input vector with weight matrix

$$0 \times 0.6342 = 0$$

[0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 1., 0., 0., 0., 0., 0., 0.]

[0.6912,	0.8765,	0.4939]	,
[0.6342,	0.7481,	0.7717]	,
[0.8395,	0.2128,	0.3696]	,
[0.4900,	0.1509,	0.0689]	,
[0.2587,	0.9171,	0.8670]	,
[0.7213,	0.9922,	0.5701]	,
[0.7598,	0.5231,	0.3666]	,
[0.5150,	0.5216,	0.9682]	,
[0.2248,	0.0261,	0.4427]	,
[0.1818,	0.6863,	0.8713]	,
[0.4192,	0.1566,	0.9004]	,
[0.8102,	0.5741,	0.4241]	,
[0.1116,	0.0466,	0.2786]	,
[0.9816,	0.7363,	0.5899]	,
[0.9224,	0.3672,	0.6972]	,
[0.1207,	0.3372,	0.2128]	,
[0.0660,	0.1524,	0.8440]	,
[0.2162,	0.5640,	0.0988]	,
[0.2605,	0.3766,	0.3502]	,
[0.2334,	0.4757,	0.7581]	,
[0.7382,	0.9807,	0.4762]	,
[0.2369,	0.8102,	0.8798]	,
[0.6932,	0.2671,	0.8018]	,
[0.9593,	0.5302,	0.4290]	,
[0.6231,	0.8825,	0.8836]	,
[0.4623,	0.8503,	0.7279]]

Multiply input vector with weight matrix

[0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 1., 0., 0., 0., 0., 0., 0., 0.]

$$1 \times 0.2605 = 0.2605$$

[0.6912,	0.8765,	0.4939],
[0.6342,	0.7481,	0.7717],
[0.8395,	0.2128,	0.3696],
[0.4900,	0.1509,	0.0689],
[0.2587,	0.9171,	0.8670],
[0.7213,	0.9922,	0.5701],
[0.7598,	0.5231,	0.3666],
[0.5150,	0.5216,	0.9682],
[0.2248,	0.0261,	0.4427],
[0.1818,	0.6863,	0.8713],
[0.4192,	0.1566,	0.9004],
[0.8102,	0.5741,	0.4241],
[0.1116,	0.0466,	0.2786],
[0.9816,	0.7363,	0.5899],
[0.9224,	0.3672,	0.6972],
[0.1207,	0.3372,	0.2128],
[0.0660,	0.1524,	0.8440],
[0.2162,	0.5640,	0.0988],
[0.2605,	0.3766,	0.3502],
[0.2334,	0.4757,	0.7581],
[0.7382,	0.9807,	0.4762],
[0.2369,	0.8102,	0.8798],
[0.6932,	0.2671,	0.8018],
[0.9593,	0.5302,	0.4290],
[0.6231,	0.8825,	0.8836],
[0.4623,	0.8503,	0.7279]

Multiply input vector with weight matrix

[0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 1., 0., 0., 0., 0., 0., 0., 0.]

$$1 \times 0.3766 = 0.3766$$

[0.6912,	0.8765,	0.4939],
[0.6342,	0.7481,	0.7717],
[0.8395,	0.2128,	0.3696],
[0.4900,	0.1509,	0.0689],
[0.2587,	0.9171,	0.8670],
[0.7213,	0.9922,	0.5701],
[0.7598,	0.5231,	0.3666],
[0.5150,	0.5216,	0.9682],
[0.2248,	0.0261,	0.4427],
[0.1818,	0.6863,	0.8713],
[0.4192,	0.1566,	0.9004],
[0.8102,	0.5741,	0.4241],
[0.1116,	0.0466,	0.2786],
[0.9816,	0.7363,	0.5899],
[0.9224,	0.3672,	0.6972],
[0.1207,	0.3372,	0.2128],
[0.0660,	0.1524,	0.8440],
[0.2162,	0.5640,	0.0988],
[0.2685,	0.3766,	0.3502],
[0.2334,	0.4757,	0.7581],
[0.7382,	0.9807,	0.4762],
[0.2369,	0.8102,	0.8798],
[0.6932,	0.2671,	0.8018],
[0.9593,	0.5302,	0.4290],
[0.6231,	0.8825,	0.8836],
[0.4623,	0.8503,	0.7279]

Multiply input vector with weight matrix

[0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 1., 0., 0., 0., 0., 0., 0., 0.]

$$1 \times 0.3502 = 0.3502$$

[0.6912,	0.8765,	0.4939],
[0.6342,	0.7481,	0.7717],
[0.8395,	0.2128,	0.3696],
[0.4900,	0.1509,	0.0689],
[0.2587,	0.9171,	0.8670],
[0.7213,	0.9922,	0.5701],
[0.7598,	0.5231,	0.3666],
[0.5150,	0.5216,	0.9682],
[0.2248,	0.0261,	0.4427],
[0.1818,	0.6863,	0.8713],
[0.4192,	0.1566,	0.9004],
[0.8102,	0.5741,	0.4241],
[0.1116,	0.0466,	0.2786],
[0.9816,	0.7363,	0.5899],
[0.9224,	0.3672,	0.6972],
[0.1207,	0.3372,	0.2128],
[0.0660,	0.1524,	0.8440],
[0.2162,	0.5640,	0.0988],
[0.2605,	0.3766,	0.3502],
[0.2334,	0.4757,	0.7581],
[0.7382,	0.9807,	0.4762],
[0.2369,	0.8102,	0.8798],
[0.6932,	0.2671,	0.8018],
[0.9593,	0.5302,	0.4290],
[0.6231,	0.8825,	0.8836],
[0.4623,	0.8503,	0.7279]

Multiply input vector with weight matrix

```
[0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 1., 0., 0., 0., 0., 0., 0., 0.]
```

Result of the vector matrix multiplication:

```
[0.2605, 0.3766, 0.3502]
```

```
[[0.6912, 0.8765, 0.4939],  
 [0.6342, 0.7481, 0.7717],  
 [0.8395, 0.2128, 0.3696],  
 [0.4900, 0.1509, 0.0689],  
 [0.2587, 0.9171, 0.8670],  
 [0.7213, 0.9922, 0.5701],  
 [0.7598, 0.5231, 0.3666],  
 [0.5150, 0.5216, 0.9682],  
 [0.2248, 0.0261, 0.4427],  
 [0.1818, 0.6863, 0.8713],  
 [0.4192, 0.1566, 0.9004],  
 [0.8102, 0.5741, 0.4241],  
 [0.1116, 0.0466, 0.2786],  
 [0.9816, 0.7363, 0.5899],  
 [0.9224, 0.3672, 0.6972],  
 [0.1207, 0.3372, 0.2128],  
 [0.0660, 0.1524, 0.8440],  
 [0.2162, 0.5640, 0.0988],  
 [0.2605, 0.3766, 0.3502],  
 [0.2334, 0.4757, 0.7581],  
 [0.7382, 0.9807, 0.4762],  
 [0.2369, 0.8102, 0.8798],  
 [0.6932, 0.2671, 0.8018],  
 [0.9593, 0.5302, 0.4290],  
 [0.6231, 0.8825, 0.8836],  
 [0.4623, 0.8503, 0.7279]]
```

**This matrix multiplication is very inefficient
compared to an index-based lookup**

Replace multiplication between one-hot encoded vector and matrix with index look-up

Result of the vector matrix multiplication:

[0.2605, 0.3766, 0.3502]

One-hot encoded (“sparse”) representation of “S U N N Y”

	A	B	C	D	E	F	G	H	I	J	K	L	M	N	O	P	Q	R	S	T	U	V	W	X	Y	Z
S	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	1	0	0	0	0	0	0	0
U	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	1	0	0	0	0	0
N	0	0	0	0	0	0	0	0	0	0	0	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0
N	0	0	0	0	0	0	0	0	0	0	0	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0
Y	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	1	0



Embedded (“dense”) representation of “S U N N Y”

[[0.9816, 0.7363, 0.5899],
[0.2605, 0.3766, 0.3502],
[0.7382, 0.9807, 0.4762],
[0.6231, 0.8825, 0.8836]]


Embedding layer

[[0.6912, 0.8765, 0.4939],
[0.6342, 0.7481, 0.7717],
[0.8395, 0.2128, 0.3696],
[0.4900, 0.1509, 0.0689],
[0.2587, 0.9171, 0.8670],
[0.7213, 0.9922, 0.5701],
[0.7598, 0.5231, 0.3666],
[0.5150, 0.5216, 0.9682],
[0.2248, 0.0261, 0.4427],
[0.1818, 0.6863, 0.8713],
[0.4192, 0.1566, 0.9004],
[0.8102, 0.5741, 0.4241],
[0.1116, 0.0466, 0.2786],
S [0.9816, 0.7363, 0.5899],
[0.9224, 0.3672, 0.6972],
[0.1207, 0.3372, 0.2128],
[0.0660, 0.1524, 0.8440],
[0.2162, 0.5640, 0.0988],
U [0.2605, 0.3766, 0.3502],
[0.2334, 0.4757, 0.7581],
N [0.7382, 0.9807, 0.4762],
[0.2369, 0.8102, 0.8798],
[0.6932, 0.2671, 0.8018],
[0.9593, 0.5302, 0.4290],
Y [0.6231, 0.8825, 0.8836],
[0.4623, 0.8503, 0.7279]]

Coming back to the character-level RNN ...

A simple RNN that predicts the next character

Desired outputs:



Sentence:

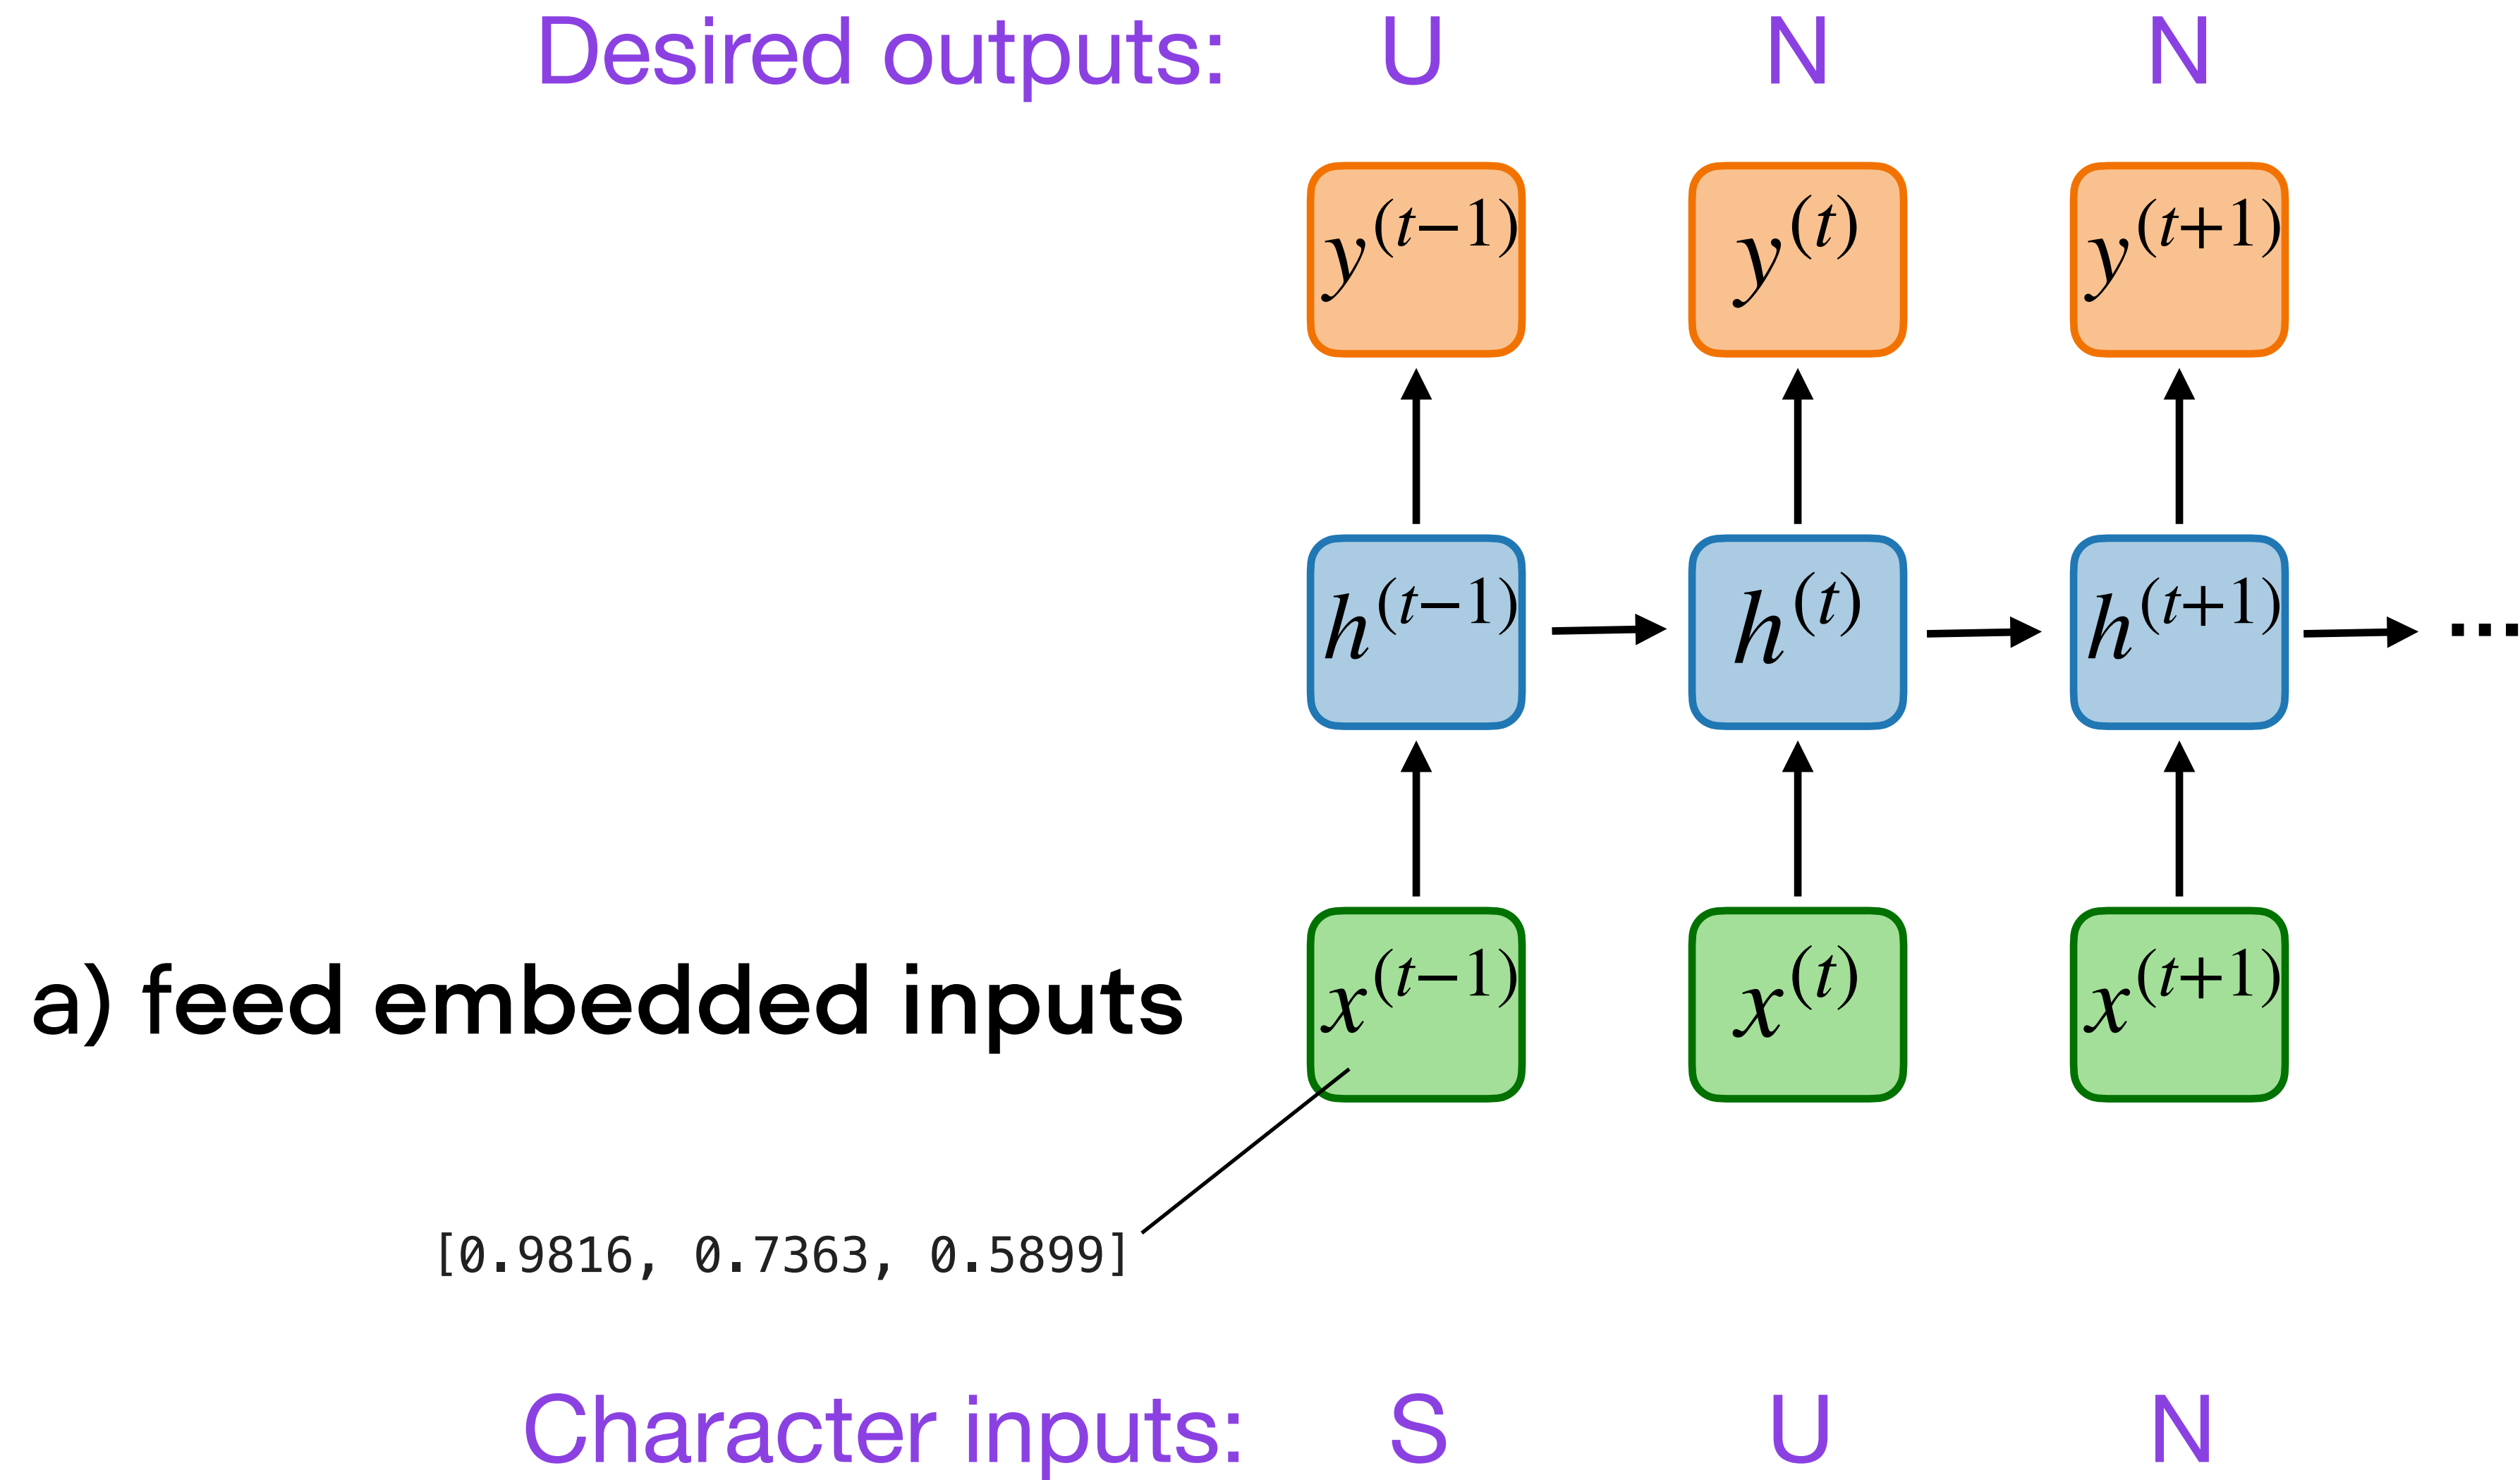
"Sunny days are the best days to go for a walk or have a picnic."

Character inputs:

S

U

N



Feed one-hot encoded inputs and let the hidden layer do the embedding.

Optimal embedding is “learned” during training.

Desired outputs:

U

N

N

$y^{(t-1)}$

$y^{(t)}$

$y^{(t+1)}$

$h^{(t-1)}$

$h^{(t)}$

$h^{(t+1)}$

...

[0.9816, 0.7363, 0.5899]

b) feed one-hot encoded inputs

$x^{(t-1)}$

$x^{(t)}$

$x^{(t+1)}$

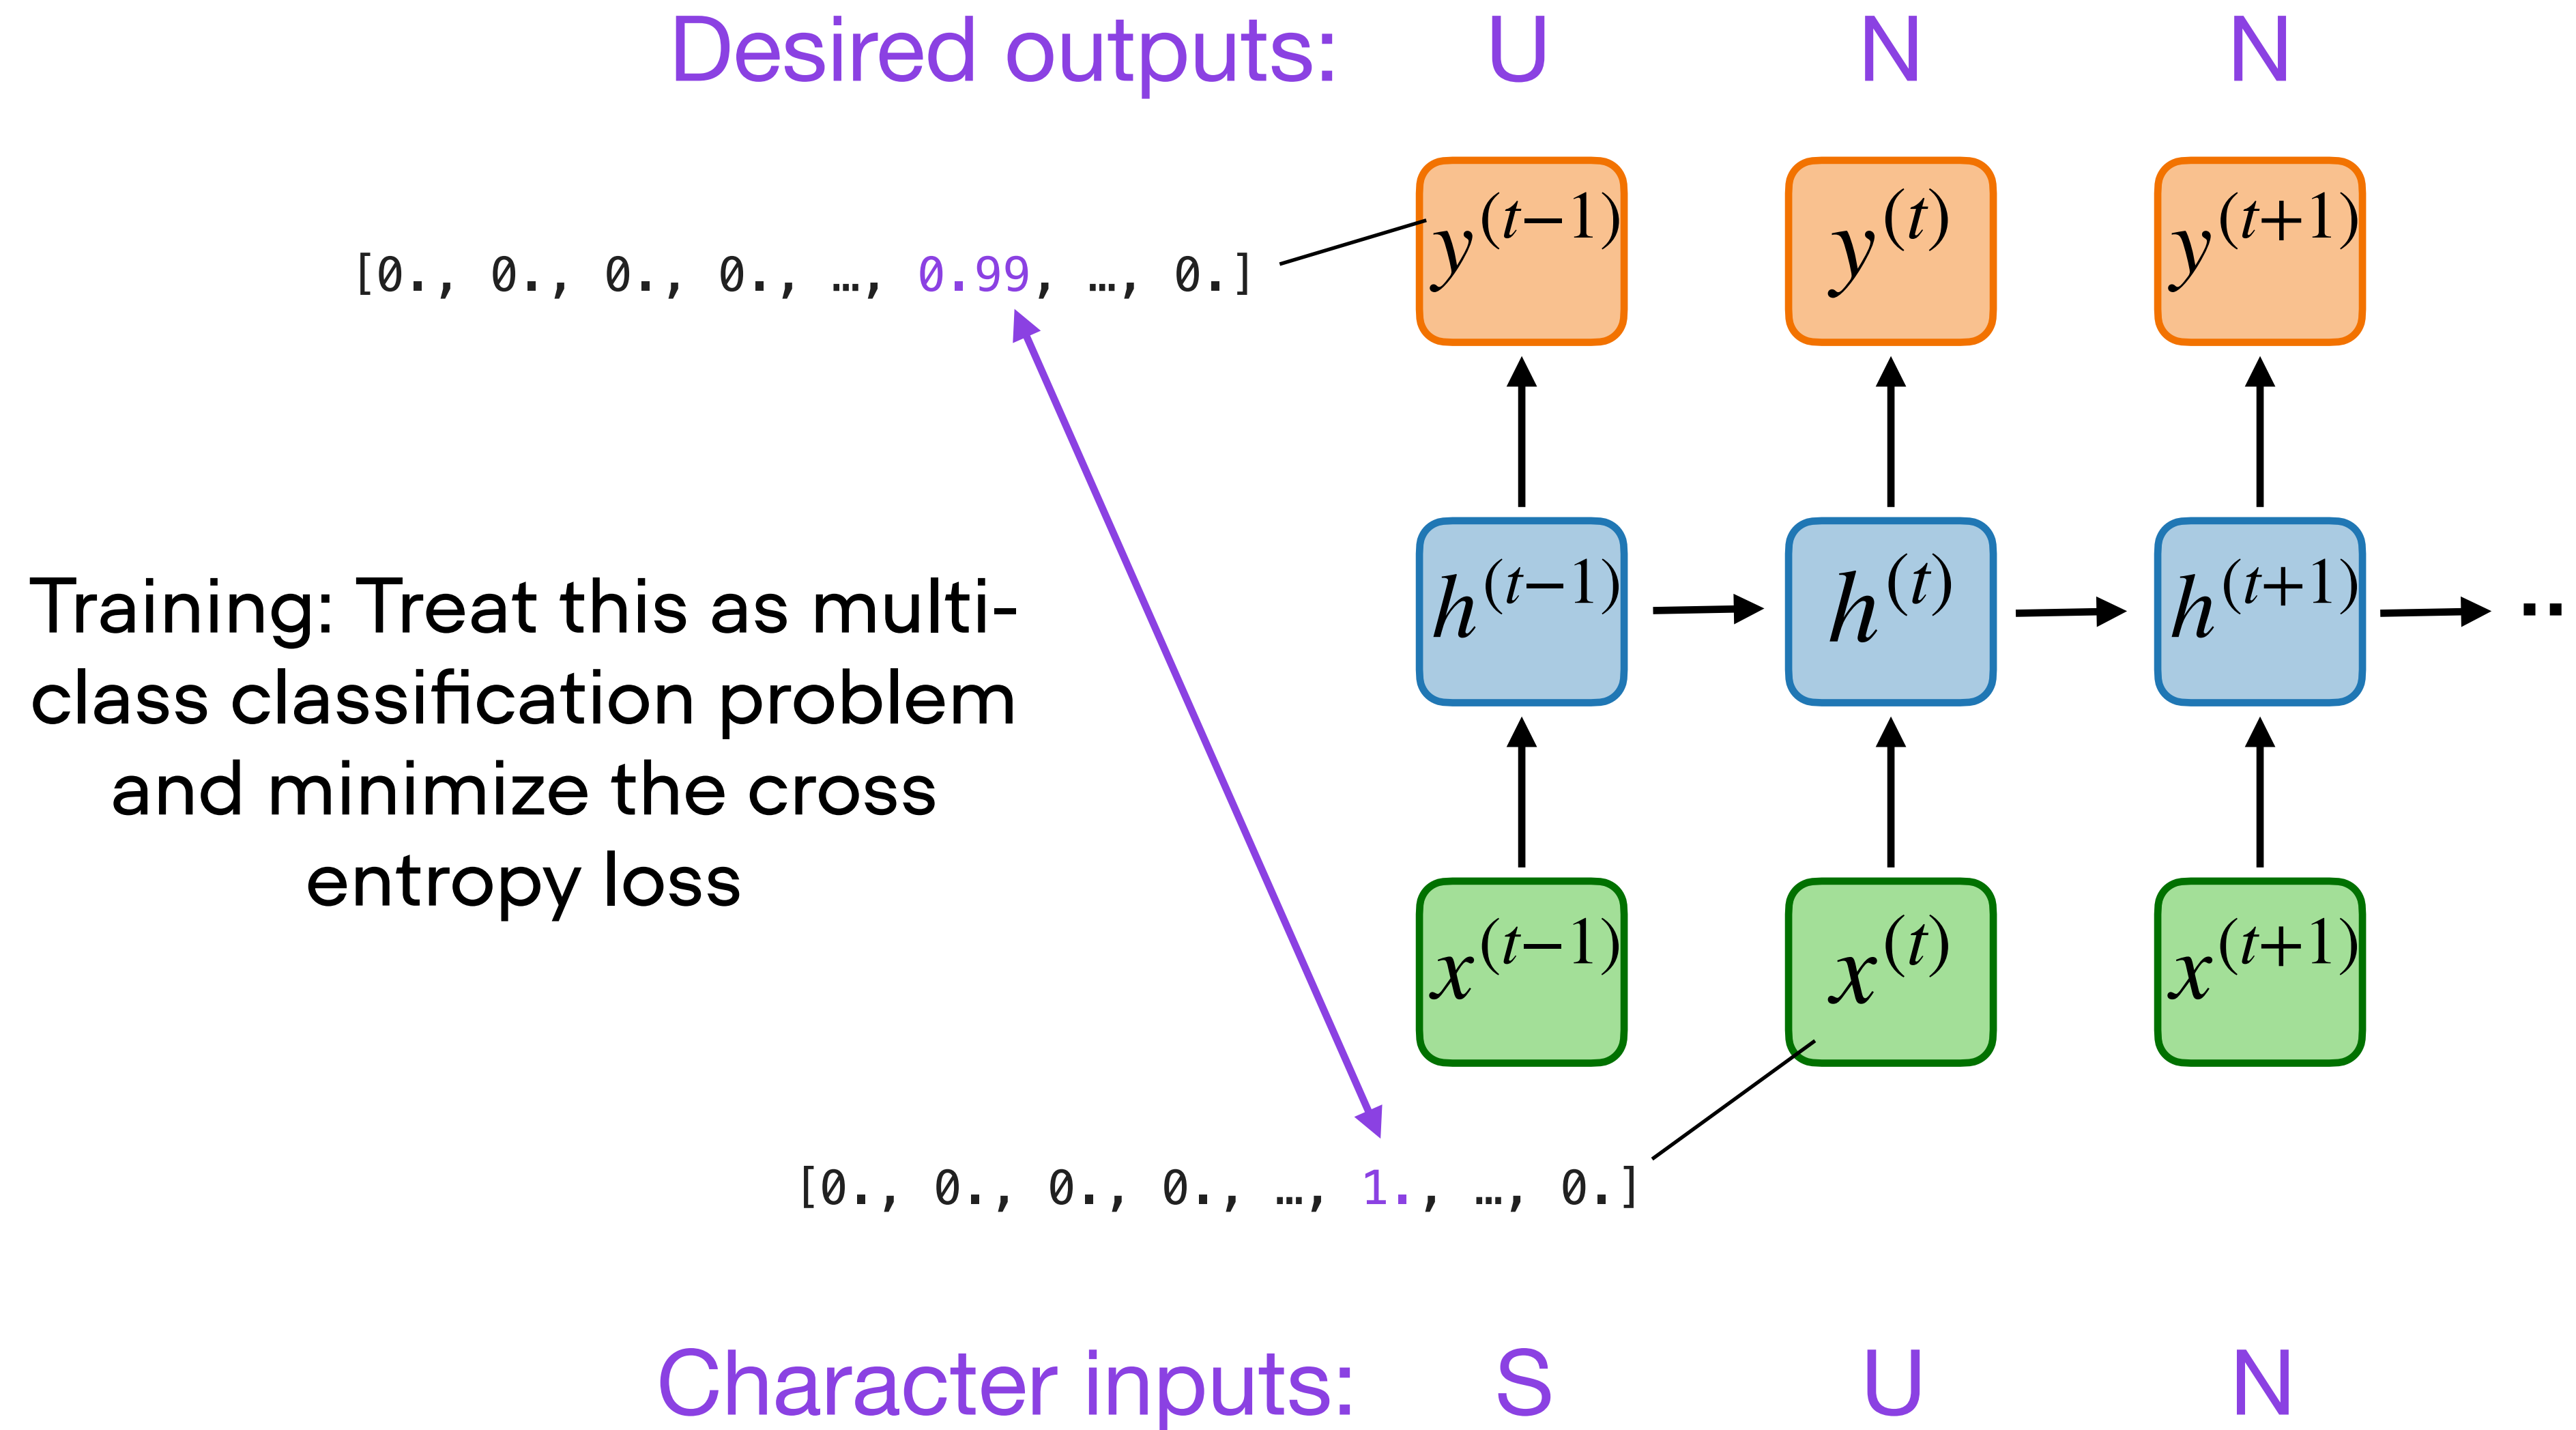
[0, 0, 0, 0, ..., 1, ..., 0]

Character inputs:

S

U

N



Word vs character embeddings

Advantages / disadvantages of character embeddings vs. word embeddings

- + require less memory (English: 26 letters + punctuation)
- + Smaller output layers
- Can create nonsense words
- Worse at capturing long-distance dependencies

Next: PyTorch Embedding Layer Example